



TÉCNICO LISBOA

Sistemas de Informação e Bases de
Dados
2020/2021

Aula 02: Introdução às Bases de Dados e ao SQL

Prof. Paulo Carreira

Class Outline

- How a database Organises Data
 - Tables, Columns, Rows, and Cells
 - Key, Primary Key, and Foreign Key
 - Table Schema and Database Schema
- Introduction to SQL
 - Creating database tables and populating them with information
 - Basic query structure: SELECT
 - Filters and Pattern Matching
 - Set operations and Joins

How a Database Organises Data

The Table

A two-dimensional grid

A **two-dimensional** grid that keeps data (facts) about
(real-world) **Entities**

ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M
78001	João Gomes Vultos Freitas	22231233	L
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L


Columns, lines, and cells

Data is organised into **columns**, **lines**, and **cells**

Columns are named and identified by **Attributes**

Lines are called 'Records' or 'Rows' →

Data is recorded in **Cells** called 'Fields' →



ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M
78001	João Gomes Vultos Freitas	22231233	L
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L

Data Types, Domains, and Formats

For each **column**, the data in **every cell** has the same:

- **Data Domain**: Non-negative, min and max value range, minimum and maximum length, allowed characters, words, or codes from a set of values
- **Data Type**: Numeric, Text, Date, Binary
- **Data Format**: 5 digits, Lowercase vs Capitalised, License Plate, YYYY/MM/DD, etc

Schema = Columns → Data Domain			
ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76555

Header

Table Data

Every Column has an associated **Domain, Type and Format**

Atomic values

- ▶ Text

```
'John Smith', 'R2D2', 'Red', ''
```

- ▶ Numeric values

```
-1, 25, +6.34, 0.5, 25e-03
```

- ▶ Floats and Doubles

```
25f, +6.34d, 0.5d
```

- ▶ Dates

```
'1999-01-08', '08-JAN-1999'
```


Table

Definition

An organisation of data into columns and lines where all lines have a similar structure, defined by the schema of the table

Algorithmically- vs. non-algorithmically definable column domains

Cannot be defined algorithmically

Can be defined algorithmically

ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S

Ordering of columns

If the **ordering of columns** is changed, information remains the same

ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M
78001	João Gomes Vultos Freitas	22231233	L
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L

Ordering of columns

In terms of information contents, the order of the columns is irrelevant

Name	ID	TAX ID	T-Shirt
João Guilherme Silva da Cunha	76440	12345678	M
Tomás Pinto dos Santos	76469	91234567	M
David Miguel Redwanz Duque	76500	89012345	L
Pedro Daniel Diz Pinela	64833	67890123	M
Guilherme de Queiróz Rebelo Brum Gomes	76464	22394856	XL
Marta Isabel de Almeida Cardoso	78230	34562732	S
Filipe Emanuel Lourenço Ramalho Fernandes	78083	82533235	L
Gabriel Filipe Queirós Mesquita Delgado Freire	78081	23134539	M
João Gomes Vultos Freitas	78001	22231233	L
Ricardo Afonso Rodrigues da Silva Oliveira	76504	56372848	L

Ordering of rows

When the **order of rows** is changed, information remains the same

ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M
78001	João Gomes Vultos Freitas	22231233	L
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L

Ordering of rows

In terms of information contents, the **order of the rows is irrelevant**

ID	Name	TAX ID	T-Shirt
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M
78001	João Gomes Vultos Freitas	22231233	L
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M

The row number is irrelevant

Keys

Key values are the way to refer **uniquely** to a row

Primary Key →

Key attributes			
ID	Name	TAX ID	T-Shirt
76440	João Guilherme Silva da Cunha	12345678	M
76469	Tomás Pinto dos Santos	91234567	M
76500	David Miguel Redwanz Duque	89012345	L
64833	Pedro Daniel Diz Pinela	67890123	M
76464	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL
78230	Marta Isabel de Almeida Cardoso	34562732	S
78082	Filipe Emanuel Lourenço Ramalho Fernandes	82522225	L

Retrieving a record given a key value is a fundamental feature/operation of database systems

76504 Ricardo Afonso Rodrigues da Silva Oliveira 56372848 L

Table Relationships

A foreign key connects two tables

Foreign Key

Employee



ID	Name	TAX ID	T-Shirt	DID
76440	João Guilherme Silva da Cunha	12345678	M	EN
76469	Tomás Pinto dos Santos	91234567	M	EN
76500	David Miguel Redwanz Duque	89012345	L	MK
64833	Pedro Daniel Diz Pinela	67890123	M	HR
76464	Guilherme de Queirós Rebelo Brum Gomes	22394856	XL	EN
78230	Marta Isabel de Almeida Cardoso	34562732	S	HR
78083	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L	EN
78081	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M	EN
78001	João Gomes Vultos Freitas	22231233	L	EN
76504	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L	MK

Department

DID	Name	Budget
HR	Human Resources	50 000
EN	Software Engineering	1 200 000
MK	Marketing	150 000

Primary Key

Master-detail related tables

SUPPLIER

Columns (Attributes, Fields)

Supplier_Number	Supplier_Name	Supplier_Street	Supplier_City	Supplier_State	Supplier_Zip
8259	CBM Inc.	74 5 th Avenue	Dayton	OH	45220
8261	B. R. Molds	1277 Gandolly Street	Cleveland	OH	49345
8263	Jackson Composites	8233 Micklin Street	Lexington	KY	56723
8444	Bryant Corporation	4315 Mill Drive	Rochester	NY	11344

Master

Rows
(Records,
Tuples)

Key Field
(Primary Key)

PART

Part_Number	Part_Name	Unit_Price	Supplier_Number
137	Door latch	22.00	8259
145	Side mirror	12.00	8444
150	Door molding	6.00	8263
152	Door lock	31.00	8259
155	Compressor	54.00	8261
178	Door handle	10.00	8259

Detail

Primary Key

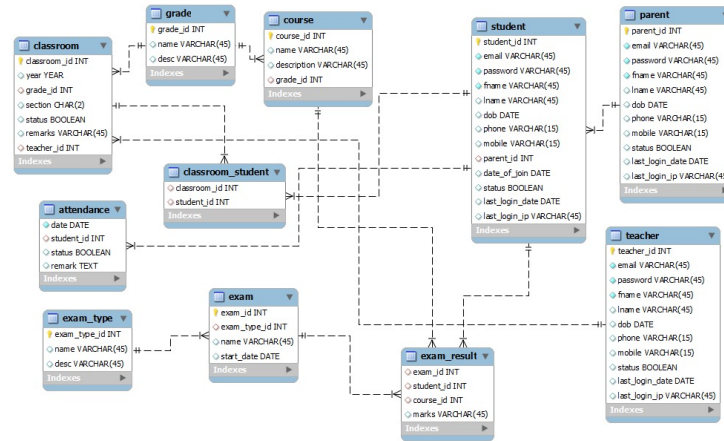
Foreign Key

Tables Concepts Summary

- **Rows:** Records data different entities
- **Columns:** Represents attributes for entity
- **Fields:** Represents attribute for entity
- **Key field:** Field used to uniquely identify each record
- **Primary key:** Field in table used for key fields
- **Foreign key:** Primary key used in second table as look-up field to identify records from original table

Databases and Database Schemas

Database Schema



1. **Individual schemas of tables:** Define which data can be captured
2. **Relationships between tables:** Define how data can be related and navigated (the navigational structure)
3. **Integrity Constraints:** Define which data are not valid in the database



What is kept in the database

Schema

The description of data

Data

The actual data contents



Database

Definition

1. An **organised collection of data** (typically of large size) that is inter-relatable
2. A **computerised system**, that allows data to be easily accessed, manipulated and updated

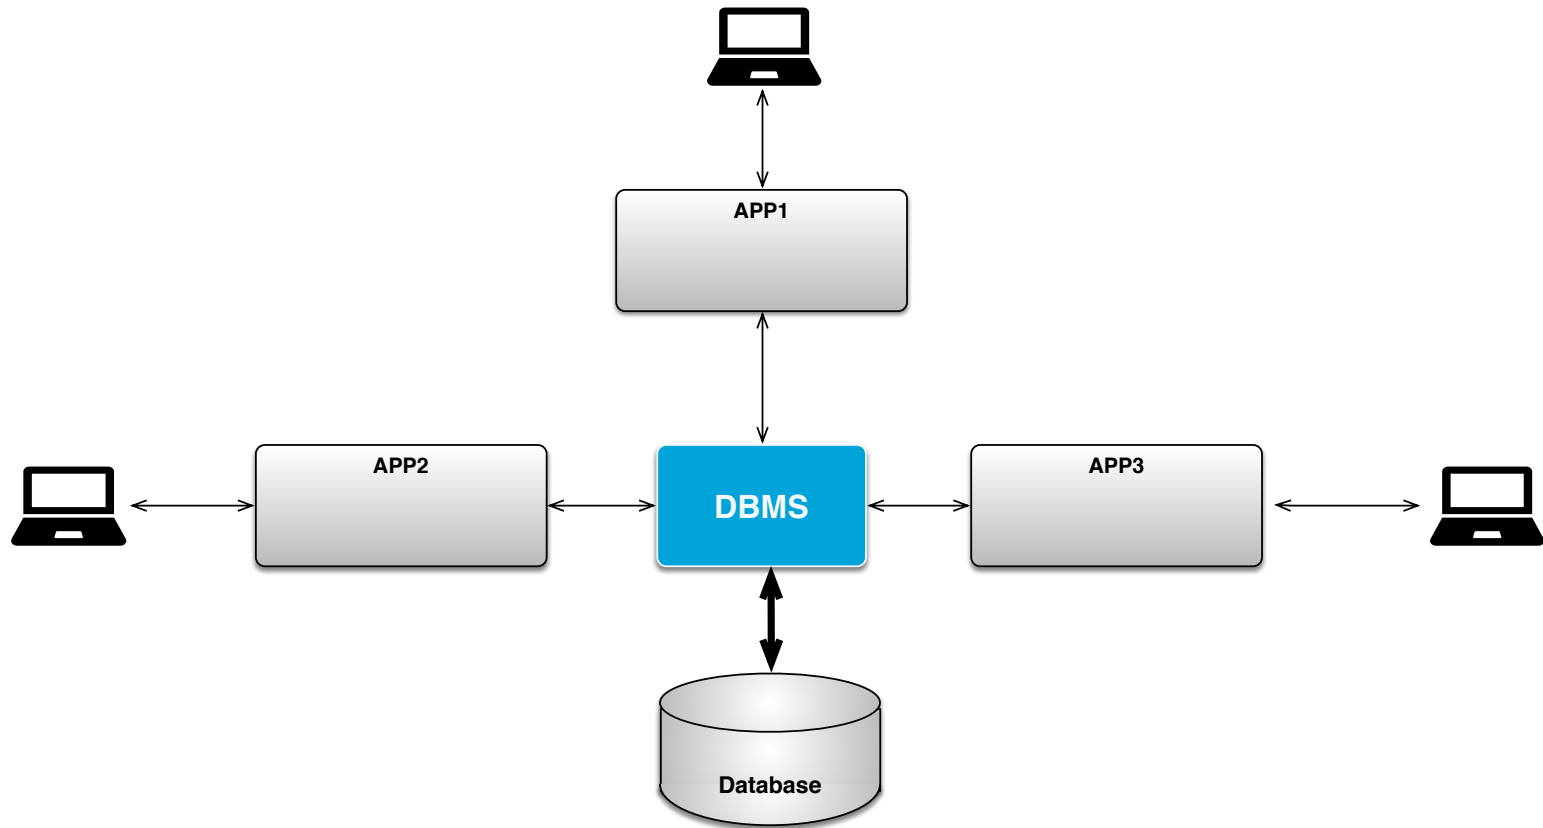
See also

What is a Database?



<https://www.youtube.com/watch?v=t8igX1f8kc4>

DBMS



Applications do not contact the database directly

Database Management System (DBMS)

Definition

A **software package** that manages a Database allows data to be easily accessed, manipulated, and updated

Introducing SQL

SQL Language

Organised into two sub-languages (parts):

1. Data Definition Language (SQL-DDL)
2. Data Query Language (SQL-DQL)

The Bank running example

Branch

branch_name	branch_city	assets
Central	Guimarães	2100000.0000
Clérigos	Oporto	3900000.0000
Downtown	Lisbon	1900000.0000
Metro	Amadora	400200.0000
Round Hill	Amadora	8000000.0000
Ship Terminal	Setúbal	400000.0000
University	Vila Real	7200000.0000
Uptown	Amadora	1700000.0000
Wine Celar	Oporto	4002800.0000

Account

account_number	branch_name	balance
A-101	Downtown	500.0000
A-102	Uptown	700.0000
A-201	Uptown	900.0000
A-215	Metro	600.0000
A-217	University	650.0000
A-222	Central	550.0000
A-305	Round Hill	800.0000
A-333	Central	750.0000
A-444	Downtown	850.0000

Depositor

customer_name	account_number
Cook	A-101
Johnson	A-101
Cook	A-102
Johnson	A-201
Brown	A-215
Iacocca	A-217
Evans	A-222
Flores	A-305
Oliver	A-333
Brown	A-444

Loan

loan_number	branch_name	amount
L-11	Round Hill	6000.0000
L-14	Downtown	4000.0000
L-15	Uptown	3000.0000
L-16	Uptown	7000.0000
L-17	Downtown	1000.0000
L-20	Downtown	8000.0000
L-21	Central	9000.0000
L-23	Central	2000.0000
L-93	Metro	5000.0000

Borrower

customer_name	loan_number
Brown	L-11
Nguyen	L-14
Cook	L-15
Iacocca	L-16
Gonzalez	L-17
Iacocca	L-17
Parker	L-20
Brown	L-21
Brown	L-23
Davis	L-93

Customer

customer_name	customer_street	customer_city
Adams	Main Street	Lisbon
Brown	Main Street	Oporto
Cook	Main Street	Lisbon
Davis	Church Street	Oporto
Evans	Forest Street	Coimbra
Flores	Station Street	Braga
Gonzalez	Sunny Street	Faro
Iacocca	Spring Steet	Coimbra
Johnson	New Street	Guimarães
King	Garden Street	Aveiro
Lopez	Grand Street	Vila Real
Martin	Royal Street	Braga
Nguyen	School Street	Castelo Branco
Oliver	First Stret	Oporto
Parker	Hope Street	Oporto



Table



Value dependency (field)

Creating Tables

The CREATE TABLE statement

A new **table** is created with the command:

```
CREATE TABLE <table-name> (  
    column1 type1,  
    ...  
    columnn typen,  
    <integrity-constraint1>,  
    ...  
    <integrity-constraintm>  
); [storage-options]
```

- ▶ The syntax of the storage options is dependent on the SGBD manufacturer

CREATE TABLE: Example

Column/field

Type

```
CREATE TABLE department(  
  did INTEGER,  
  name VARCHAR(80),  
  budget NUMERIC(12,4),  
  PRIMARY KEY(did));
```

Integrity Constraint

Removing a table

- ▶ Removing a table

```
DROP TABLE <table-name>
```

- ▶ Removing the table **department** from the database

```
DROP TABLE department;
```

Specifying foreign key constraints

```
CREATE TABLE account
(account_number CHAR(5),
 branch_name VARCHAR(80),
 balance NUMERIC(16,4),
 PRIMARY KEY(account_number),
 FOREIGN KEY(branch_name)
 REFERENCES branch(branch_name));
```

Creating, updating, and deleting rows

Inserting and querying rows

```
INSERT INTO <table-name>  
VALUES(<field-values>)
```

- ▶ Inserting a record into a table

```
INSERT INTO department VALUES(1,  
    'Finance', 1000.0);
```

- ▶ Querying the records of a table

```
SELECT * FROM <table-name>
```

Schema Definition – Example 1

```
CREATE TABLE department(  
    did INTEGER,  
    name VARCHAR(80),  
    budget NUMERIC(12,2));
```

```
INSERT INTO department VALUES(1, 'Finance', 1000.0);  
INSERT INTO department VALUES(2, 'Marketing', 2000.0);
```

```
SELECT * FROM department;
```

did		name		budget
-----	+	-----	+	-----
1		Finance		1000.0000
2		Marketing		2000.0000

Schema Definition – Example 2

```
INSERT INTO department VALUES(1, 'Finance', 1000.0);  
INSERT INTO department VALUES(2, 'Marketing', 2000.0);
```

```
INSERT INTO department VALUES(1, 'Engineering', 5000.0);  
INSERT INTO department VALUES(3, 'Marketing', 3000.0);
```

```
SELECT * FROM department;
```

did	name	budget
1	Finance	1000.0000
2	Marketing	2000.0000
1	Engineering	5000.0000
3	Marketing	3000.0000

⚠ We have a repeated ***did*** key and a repeated ***name***

Schema Definition – Example 3

```
CREATE TABLE department(  
    did INTEGER,  
    name VARCHAR(80),  
    budget NUMERIC(12,4),  
    PRIMARY KEY(did),  
    UNIQUE(name));
```

```
INSERT INTO department VALUES(1, 'Finance', 1000.0);  
INSERT INTO department VALUES(2, 'Marketing', 2000.0);
```

```
INSERT INTO department VALUES(1, 'Engineering', 5000.0);  
INSERT INTO department VALUES(3, 'Marketing', 3000.0);
```

 The last two statements fail!

Retrieving Data

What is the result of a query?

sid	name	grade
53666	Maria	10
53688	Joseph	11
53650	Chris	12
53831	Ana	13
53832	John	14

Who are the students with grades above 11?

sid	name	grade
53650	Chris	12
53831	Ana	13
53832	John	14

What is the percent of students that finished with a grade greater than 13?

percent
20%

Queries convert data into information!

Basic Query Structure

- ▶ SQL is based on set and relational operations with certain modifications and enhancements
- ▶ The fundamental idea is that of selecting the records that satisfy a given set of criteria; SQL allows for the definition of arbitrarily complex criteria.
- ▶ A typical SQL query has the form:

```
select C1, C2, ..., Cn  
from table  
where condition
```

- ▶ The result of an SQL query is also a **table**

The **SELECT** clause

- ▶ The select clause lists the attributes desired in the result of a query
- ▶ **Example:** find the names of all branches in the loan relation:

```
select branch_name  
from loan;
```

```
branch_name  
-----  
Downtown  
Central  
Uptown  
Downtown  
Metro  
Round Hill  
Uptown  
Downtown  
Central
```

NB: SQL names are case insensitive (i.e., you may use upper- or lower-case characters)

The **SELECT** clause (cont.)

- ▶ An asterisk in the select clause denotes “*all attributes*”

```
select *  
from loan;
```

loan_number	branch_name	amount
L-17	Downtown	1000.0000
L-23	Central	2000.0000
L-15	Uptown	3000.0000
L-14	Downtown	4000.0000
L-93	Metro	5000.0000
L-11	Round Hill	6000.0000
L-16	Uptown	7000.0000
L-20	Downtown	8000.0000
L-21	Central	9000.0000

Expressions and derived fields

The **SELECT** clause (cont.)

- ▶ The **select** clause can contain arithmetic expressions involving the operations $+$, $-$, $*$, and $/$, as well as functions operating on constants or attributes of tuples

```
select loan_number, branch_name, amount * 1.1  
from loan
```

- ▶ Returns a relation that is the same as the loan relation, except that the value of the attribute amount is multiplied by 1.1

Field and expression aliases

- ▶ SQL allows renaming tables and columns using the **AS** clause:

old-name **AS** new-name

- ▶ **Example:** Find the new balance of every account after subtracting 10€ in expenses

```
select account_number as acct, (balance-10) as new_bal
from account;
```

acct		new_bal
A-101		490.0000
A-215		590.0000
A-102		690.0000
A-305		790.0000
A-201		890.0000
A-222		540.0000
A-217		640.0000
A-333		740.0000
A-444		840.0000

Filtering rows

The **WHERE** clause

- ▶ The **where** clause specifies conditions that the table rows must satisfy (i.e., acts as a **filter**)
- ▶ **Example:** Find all loan numbers of loans made at the 'Central' branch with loan amounts greater than 5000€

```
select loan_number, amount
from loan
where branch_name = 'Central'
      and amount > 5000
```

loan_number	amount
L-21	9000.0000

The **WHERE** clause

- ▶ Comparison results can be combined using the logical connectives **and**, **or**, and **not**
- ▶ Comparisons can be applied to results of arithmetic expressions

The **WHERE** clause

- ▶ SQL includes the **between ... and** operator
- ▶ **Example:** Find the loan number of those loans whose amounts are between 5 000€ and €10 000 (i.e., $\geq \text{€}5\,000$ and $\leq \text{€}10\,000$)

```
select loan_number
from loan
where amount between 5000 and 10000
```

```
loan_number
-----
L-93
L-11
L-16
L-20
L-21
```

Handling Duplicates

The **SELECT DISTINCT** clause

- ▶ SQL allows duplicates in relations as well as in query results
- ▶ To force the elimination of duplicates, insert the keyword **distinct**
- ▶ **Example:** find the names of all branches in the loan relations, **and remove the duplicates**

```
SELECT DISTINCT branch_name  
FROM loan;
```

branch_name

Central

Uptown

Round Hill

Downtown

Metro

Searching and Pattern Matching

String Pattern Matching

- ▶ The operator **LIKE** uses patterns that are described using two special characters:
- ▶ **Underscore** '_': matches any (at least one) single character.
- ▶ **Percent** '%': matches any sequence of variable length (including the empty sequence)

Pattern	Matches
Ana%	Ana, Anastasia, Anacleto, Anafora, Ananás, ...
%ana%	Banana, Cana, Panacota, Feliciano, ...
Ana_%s	Ana Santos, Ana Martins, Ananás
Ana %s	Ana Santos, Ana Martins, ...
__-__-__	XT-09-10, 20-45-XY, 10-AD-90, AB-CD-EF, ...

String Pattern Matching

- ▶ **Example:** Find the street of all customers whose name starts with “Ana”

```
select customer_name, customer_street  
from customer  
where customer_name like 'Ana%'
```

Ana, Anastasia, Anacleto, Anafora, Ananás, ...

- ▶ **Example:** Find the street of all customers whose name has “ana”

```
select customer_name, customer_street  
from customer  
where customer_name like '%ana%'
```

Banana, Cana, Panacota, Feliciano, ...

Exercise

Exercise

- A. Get the data of all accounts
- B. Get the number of all accounts with a balance between 700€ and 900€
- C. Get the name of all customers whose street name starts in 'C'

```
SELECT *  
FROM account
```

```
SELECT account_number  
FROM account  
WHERE balance BETWEEN 700 AND 900
```

```
SELECT customer_name  
FROM customer  
WHERE customer_street LIKE 'C%'
```

Sorting Rows

Changing the display order of rows

- ▶ List in alphabetic order the names of all customers

```
SELECT DISTINCT customer_name
FROM customer
ORDER BY customer_name;
```

- ▶ We may specify **desc** for descending order or **asc** for ascending order (the default)

```
SELECT DISTINCT customer_name
FROM customer
ORDER BY customer_name DESC
```

customer_name

Adams
Brown
Cook
Davis
Evans
Flores
Gonzalez
Iacocca
Johnson
King
Lopez
Martin
Nguyen
Oliver
Parker

customer_name

Parker
Oliver
Nguyen
Martin
Lopez
King
Johnson
Iacocca
Gonzalez
Flores
Evans
Davis
Cook
Brown
Adams

Order in DBMS

► Order-dependent Storage

- Each application may require data sorted in a different way; even distinct screens in the same application
- If records were stored ordered by a certain column, then sorting them by other order is costly

► Order-independent Storage

- The order by which records are organised in the storage should be irrelevant (data are as a principle independent from the application)
- DBMSs provide efficient ways to sort records in whatever way the application requires

Set Operations

Set Operations

Union of tables

```
SELECT ...  
UNION  
SELECT ...
```

A	B	C

A	B	C

A	B	C

Intersection of tables

```
SELECT ...  
intersect  
SELECT ...
```

A	B	C

A	B	C

A	B	C

Difference of tables

```
SELECT ...  
except  
SELECT ...
```

A	B	C

A	B	C

A	B	C

⚠ Set operators eliminate duplicate tuples by default. The **all** clause should be specified in order to consider the duplicates.

Set Operations

- ▶ Find all customers who have a loan, an account, or both

```
(SELECT customer_name FROM depositor)  
UNION  
(SELECT customer_name FROM borrower)
```

- ▶ Find all customers who have both a loan and an account

```
(SELECT customer_name FROM depositor)  
INTERSECT  
(SELECT customer_name FROM borrower)
```

- ▶ Find all customers who have an account but no loan

```
(SELECT customer_name FROM depositor)  
EXCEPT  
(SELECT customer_name FROM borrower)
```


Navigating by Joining Relations

JOIN Example

List the names of all *employees* along their corresponding *departments*

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X	●————●	X	Marketing	Damaia
2	Barbara	Y	●————●	Y	Sales	Amadora
3	Carlos	Z	●————●	Z	Production	Buraca
4	Duarte	W	●————●	W	Support	Cacém

```
SELECT e.ename, d.name
FROM Employee e JOIN Department d
ON e.did = d.did
```

e.ename	d.name
Alice	Marketing
Barbara	Sales
Carlos	Production
Duarte	Support

JOIN Operations

Join Operations

Joining two tables

A	B	C

C	D	E

Join Operations

Results on a table that ties together the records that agree on (that match) the values of a column (or columns) subject to a condition

A	B	C	D	E

Like extending one table with another

Behaviour of INNER JOIN

Record Mapping

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X	● — ●	X	Marketing	Damaia
2	Barbara	Y	● — ●	Y	Sales	Amadora
3	Carlos	Z	● — ●	Z	Production	Buraca
4	Duarte	W	● — ●	W	Support	Cacém

```
SELECT *  
FROM Employee e JOIN Department d  
ON e.did = d.did
```

e.eid	e.name	did	d.name	e.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
3	Carlos	Z	Production	Buraca
4	Duarte	W	Support	Cacém

Values are
all distinct

One-to-One = Mapping

Record Multiplication

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X		X	Marketing	Damaia
2	Barbara	X		Y	Sales	Amadora
3	Carlos	Z		Z	Production	Buraca
4	Duarte	W		W	Support	Cacém

```
SELECT *
FROM Employee e JOIN Department d
ON e.did = d.did
```

Some values
repeat

e.ei	ename	did	dname	d.loc
1	Alice	X	Marketing	Damaia
2	Barbara	X	Marketing	Damaia
3	Carlos	Z	Production	Buraca
4	Duarte	W	Support	Cacém

Many-to-One = Multiplication

Record Filtering

Employee			Department		
eid	ename	did	did	dname	loc
1	Alice	X	X	Marketing	Damaia
2	Barbara	Y	Y	Sales	Amadora
3	Carlos	T	W	Support	Cacém
4	Duarte	W			

```
SELECT *  
FROM Employee e JOIN Department d  
ON e.did = d.did
```

Some
values
don't have
a match

e.eid	e.ename	did	d.dname	e.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
4	Duarte	W	Production	Buraca

Zero-to-One = Filter

Behaviour of OUTTER JOIN

LEFT OUTER JOIN

Employee			Department		
eid	name	did	did	name	loc
1	Alice	X	X	Marketing	Damaia
2	Barbara	Y	Y	Sales	Amadora
3	Carlos	T	Z	Production	Buraca
4	Duarte	W	W	Support	Cacém

```
SELECT e.name, d.name
FROM Employee e LEFT OUTER JOIN Department d
ON d.did = d.did
```

e.ei	e.name	did	d.name	d.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
3	Carlos	T	NULL	NULL
4	Duarte	W	Support	Buraca

Blank line
with NULL
values

One output row for every row of the table on the **LEFT**

RIGHT OUTER JOIN

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X	●—●	X	Marketing	Damaia
2	Barbara	Y	●—●	Y	Sales	Amadora
3	Carlos	T	●	Z	Production	Buraca
4	Duarte	W	●—●	W	Support	Cacém

```
SELECT e.name, d.name
FROM Employee e RIGHT OUTER JOIN Department d
ON d.did = d.did
```

e.ei	e.name	did	d.name	e.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
NULL	NULL	Z	Production	Buraca
4	Duarte	W	Support	Buraca

Blank line
with NULL
values

One output row for every row of the table on the **RIGHT**

FULL OUTER JOIN

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X	●—●	X	Marketing	Damaia
2	Barbara	Y	●—●	Y	Sales	Amadora
3	Carlos	T	● ●	Z	Production	Buraca
4	Duarte	W	●—●	W	Support	Cacém

```
SELECT e.name, d.name
FROM Employee e FULL OUTER JOIN Department d
ON d.did = d.did
```

e.eid	e.name	did	d.name	e.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
3	Carlos	T	NULL	NULL
	NULL	Z	Production	Buraca
4	Duarte	W	Support	Cacém

One output row for every row on **EACH** table


Querying the Bank Database with a Join

Example

Find the names of the *depositors* with *accounts* with more than 750 € in balance

Step 1: Find all Accounts with balance > 750€


```
SELECT *  
FROM account  
WHERE balance > 750  
ORDER BY balance;
```



account_number	branch_name	balance
A-305	Round Hill	800.0000
A-444	Downtown	850.0000
A-201	Uptown	900.0000

Step 2: Find the depositors of those accounts

```
SELECT *  
FROM depositor  
ORDER BY account_number;
```



customer_name	account_number
Cook	A-101
Johnson	A-101
Cook	A-102
Johnson	A-201
Brown	A-215
Iacocca	A-217
Evans	A-222
Flores	A-305
Oliver	A-333
Brown	A-444

Example

*Find the **names of the customers** with accounts having more than 750 € in balance*

Joining Account and Depositor on the attribute 'account_number'

```
SELECT *  
FROM account a JOIN depositor d  
  ON a.account_number = d.account_number  
WHERE balance > 750;
```


We are selecting all attributes

account_number	branch_name	balance	customer_name
A-305	Round Hill	800.0000	Flores
A-201	Uptown	900.0000	Johnson
A-444	Downtown	850.0000	Brown

Example

*Find the **names** of the customers with accounts having more than 750 € in balance*

```
SELECT customer_name
FROM account a JOIN depositor d
  ON a.account_number = d.account_number
WHERE balance > 750;
```



customer_name

Flores
Johnson
Brown


Still joining tables but
only selection the
customer_name

Cartesian Product

Cartesian Product

Employee

eid	name	did
1	Alice	X
2	Barbara	Y
3	Carlos	Z



Department

did	name	loc
X	Marketing	Damaia
Y	Sales	Amadora
Z	Production	Buraca

```
SELECT e.name, d.name
FROM Employee e, Department d
```


e.ei	e.name	e.did	d.di	d.name	e.loc
1	Alice	X	X	Marketing	Damaia
2	Barbara	Y	X	Marketing	Damaia
3	Carlos	T	X	Marketing	Damaia
1	Alice	X	Y	Sales	Amadora
2	Barbara	Y	Y	Sales	Amadora
3	Carlos	T	Y	Sales	Amadora
1	Alice	X	Z	Production	Buraca
2	Barbara	Y	Z	Production	Buraca
3	Carlos	T	Z	Production	Buraca

Employee X Department

The FROM clause lists the tables involved in the query

Outputs all combinations of rows, irrespective of the contents

Cartesian Product with a Filter

Employee				Department		
eid	name	did		did	name	loc
1	Alice	X		X	Marketing	Damaia
2	Barbara	Y		Y	Sales	Amadora
3	Carlos	Z		Z	Production	Buraca

```
SELECT e.name, d.name
FROM Employee e, Department d
WHERE e.did = d.did
```

Query has
a filter

e.eid	e.name	did	d.name	e.loc
1	Alice	X	Marketing	Damaia
2	Barbara	Y	Sales	Amadora
3	Carlos	Z	Production	Buraca

Rows that **do not match** the filter are dropped

```
SELECT *
FROM Employee e JOIN Department d
ON e.did = d.did
```

The Cartesian product

depositor × account

```
SELECT *  
FROM depositor d, account a
```

customer_name	account_number	account_number	branch_name	balance
Johnson	A-101	A-101	Downtown	500.0000
Johnson	A-101	A-215	Metro	600.0000
Johnson	A-101	A-102	Uptown	700.0000
Johnson	A-101	A-305	Round Hill	800.0000
Johnson	A-101	A-201	Uptown	900.0000
Johnson	A-101	A-222	Central	550.0000
Johnson	A-101	A-217	University	650.0000
Johnson	A-101	A-333	Central	750.0000
Johnson	A-101	A-444	Downtown	850.0000
Brown	A-215	A-101	Downtown	500.0000
Brown	A-215	A-215	Metro	600.0000
Brown	A-215	A-102	Uptown	700.0000
Brown	A-215	A-305	Round Hill	800.0000
Brown				900.0000
Brown				550.0000
Brown				650.0000
Brown				750.0000
Brown	A-215	A-444	Downtown	850.0000

⚠ What happens if we filter the rows on
d.account_number = a.account number?

Cartesian product with a filter

- What happens when we filter a Cartesian product?

```
SELECT *  
FROM depositor d, account a  
WHERE d.account_number = a.account_number;
```


customer_name	account_number	account_number	branch_name	balance
Johnson	A-101	A-101	Downtown	500.0000
Brown	A-215	A-215	Metro	600.0000
Cook	A-102	A-102	Uptown	700.0000
Cook	A-101	A-101	Downtown	500.0000
Flores	A-305	A-305	Round Hill	800.0000
Johnson	A-201	A-201	Uptown	900.0000
Iacocca	A-217	A-217	University	650.0000
EY				550.0000
OT				750.0000
BR				850.0000

```
SELECT *  
FROM account a JOIN depositor d  
ON a.account_number = d.account_number
```

Example

Find the names of the cities of the customers with accounts having more than 750 € in balance

```
SELECT customer_city
FROM account a, depositor d, customer c
WHERE a.account_number = d.account_number
      AND c.customer_name = d.customer_name
      AND balance > 750;
```



customer_city
Oporto
Braga
Cascais