



TÉCNICO LISBOA

Sistemas de Informação e Bases de Dados 2020/2021

Class 07: Translating E-A to SQL

Prof. Paulo Carreira






Class Outline

- ☐ Motivation
- ☐ Revisiting Referential Integrity
- ☐ Translating Entities and Attributes
- ☐ SQL Data Types
- ☐ Translating Associations

Motivation

Translation Example



| Employee | | |
|----------|----------|------------|
| eid | name | birthdate |
| 1 | Alice | 10/10/1995 |
| 2 | Bob | 03/02/1996 |
| 3 | Caroline | 04/04/1997 |
| 4 | Daniel | 03/04/1998 |
| 5 | Eduard | 10/03/1994 |

| works_at | |
|----------|-----|
| eid | did |
| 1 | 1 |
| 2 | 2 |
| 3 | 3 |
| 4 | 3 |
| 5 | 3 |

| Department | | |
|------------|-----------|----------|
| did | name | location |
| 1 | Finance | Buraca |
| 2 | Marketing | Damaia |
| 3 | Sales | Chelas |

⚠ **works_at** encodes the association of Employees to Departments
All VALUES must be coherent

Relations and Attributes

```
CREATE TABLE employee(  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    bdate DATE NOT NULL,  
    PRIMARY KEY (eid)  
);
```

```
CREATE TABLE department(  
    did INTEGER,  
    name VARCHAR(20) NOT NULL,  
    location VARCHAR(20) NOT NULL,  
    PRIMARY KEY (did)  
);
```

```
INSERT INTO employee VALUES(1, 'Alice', '1995-10-10');  
INSERT INTO employee VALUES(2, 'Bob', '1996-03-02');
```

```
INSERT INTO department VALUES(1, 'Finance', 'Buraca');  
INSERT INTO department VALUES(2, 'Marketing', 'Damaia');
```

Relations and Attributes

```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    bdate DATE NOT NULL,  
    PRIMARY KEY(eid)  
);
```

```
CREATE TABLE department (  
    did INTEGER,  
    name VARCHAR(20) NOT NULL,  
    location VARCHAR(20) NOT NULL,  
    PRIMARY KEY(did)  
);
```

```
CREATE TABLE works_at(  
    eid INTEGER,  
    did INTEGER,  
    PRIMARY KEY (eid, did)  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

Prevents inputting invalid *eid* or *did* values

```
INSERT INTO works_at VALUES(1, 1);  
INSERT INTO works_at VALUES(2, 1);  
INSERT INTO works_at VALUES(2, 99);
```

Referential Integrity

Table Relationships

A foreign key connects two tables

Foreign Key

Employee



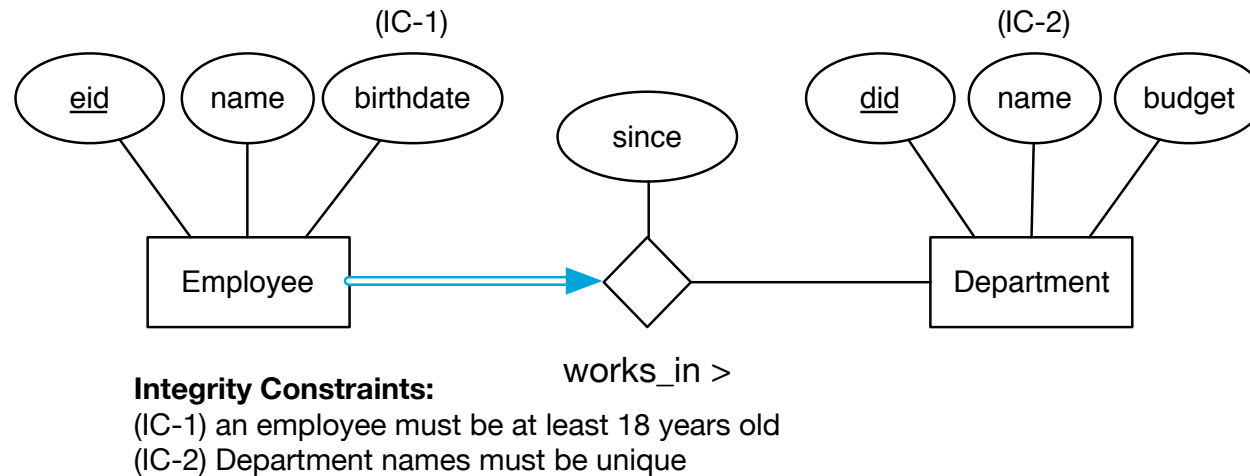
| ID | Name | Tax id | T-Shirt | DID |
|-----|--|----------|---------|-----|
| 001 | João Guilherme Silva da Cunha | 12345678 | M | EN |
| 002 | Tomás Pinto dos Santos | 91234567 | M | EN |
| 003 | David Miguel Redwanz Duque | 89012345 | L | MK |
| 004 | Pedro Daniel Diz Pinela | 67890123 | M | HR |
| 005 | Guilherme de Queiróz Rebelo Brum Gomes | 22394856 | XL | EN |
| 006 | Marta Isabel de Almeida Cardoso | 34562732 | S | HR |
| 007 | Filipe Emanuel Lourenço Ramalho Fernandes | 82533235 | L | EN |
| 008 | Gabriel Filipe Queirós Mesquita Delgado Freire | 23134539 | M | EN |
| 009 | João Gomes Vultos Freitas | 22231233 | L | EN |
| 010 | Ricardo Afonso Rodrigues da Silva Oliveira | 56372848 | L | MK |

Department

| DID | Name | Budget |
|-----|----------------------|-----------|
| HR | Human Resources | 50 000 |
| EN | Software Engineering | 1 200 000 |
| MK | Marketing | 150 000 |

↑ Primary Key

How to translate this case?



⚠ All VALUES in **Employee.dep** must exist as PRIMARY KEY the **Department.did**


Employee

| eid | name | birthdat | since | dep |
|-----|----------|------------|------------|-----|
| 1 | Alice | 10/10/1995 | 10/10/2015 | 1 |
| 2 | Bob | 03/02/1996 | 01/04/2014 | 2 |
| 3 | Caroline | 04/04/1997 | 15/04/2015 | 3 |
| 4 | Daniel | 03/04/1998 | 01/03/2018 | 3 |
| 5 | Eduard | 10/03/1994 | 01/01/2019 | 3 |

Department

| did | name | location |
|-----|-----------|----------|
| 1 | Finance | Buraca |
| 2 | Chemistry | Damaia |
| 3 | Sales | Chelas |

Operations that violate Referential Integrity




- **Removing** lines from *Department*: We cannot remove *departments* to which *employees* are still associated
- **Updating VALUES** on *Department*: We cannot change VALUES on *department* that imply changing VALUES on *employee*
- **Inserting** lines on *Employee*: We cannot add *employees* on *departments* that do not exist

Referential Integrity Constraints (or Foreign Keys)

- The most common constraint involving **two tables** is the **Referential Integrity** constraint
- Data in one table **must be always coherent with the data of** another table. A table is usually related to other tables.

Translating Entities and Attributes

Entities



```
CREATE TABLE employee (  
    eid    INTEGER,  
    name   VARCHAR(80) NOT NULL,  
    bdate  DATE NOT NULL,  
    PRIMARY KEY(eid)  
);
```

1. **Entities** result in a **table** with corresponding attributes
2. **PRIMARY KEY** is the same
3. **Constraints** have also to be translated (will see how later on...)

Data Types and Values

datatype Families

Text Types

Varchar

Char

Text

'John Smith', 'R2D2', 'Red', ''

Numeric Types

Character strings

Integer

Fixed Point

Floating Point

-1, 25, +6.34, 0.5, 25e-03

Date & Time

Date

Time

Timestamp

'2029-01-01', '08-JAN-2029 10:35:02'

Text Types

VARCHAR(n)

Variable length
character string
max size n

$n < 4000$

CHAR(n)

Character string
with fixed size n

$n < 4000$

TEXT

Variable length
text (multi-line)
field (typically
limited to 65535
characters)

*length typically
limited to 65535
characters*

Integer Numeric Types

INTEGER

An integer with
up to 9 digits

4 bytes

-2^{31}
-2 147 483 648
to
 $+2^{31}-1$
2 147 483 647

SMALLINT

An integer with
up to 4 digits

2 bytes

-2^{15}
-32 768
to
 $+2^{15}-1$
32 767

BIGINT

An integer with
up to 18 digits

8 bytes

-2^{63}
-9 223 372 036
854 775 808
to
 $+2^{63}-1$
9 223 372 036
854 775 807

Fixed Point Number Types

NUMERIC(p, s)

- A numeric value with arbitrary precision
 - **Precision p** : total number of digits, must be positive ($p > 0$)
- **Scale s** : number of digits to right of the decimal point, can be zero but must always be smaller than the precision ($0 < s < p$)
 - Up to 131072 digits before the decimal point
 - Up to 16383 digits after the decimal point

Whenever s is zero, we can write **NUMERIC(p)**

Floating Point Numeric

| REAL | DOUBLE |
|--|---|
| variable-precision, inexact 6 digits | variable-precision, inexact 15 digits |
| 4 bytes | 8 bytes |
| 1E-37 to 1E+37 | 1E-307 to 1E+308 |

Fixed- vs. Floating Point

FIXED POINT

FLOATING POINT

| | | |
|---------------|---|--|
| Precision | Fixed precision | Variable-precision |
| Storage | Stores numbers exactly | Stores numbers inexactly (as approximations) |
| Retrieval | Retrieved values never show any discrepancies | Retrieved values may show discrepancies from values stored |
| Speed | Slower calculations (*) | Faster calculations |
| Money Amounts | Can be safely used for monetary values | Should never be used for monetary values |

(*) In practice, this difference is often neglectable

Behaviour of Fixed and Float

Behaviour of Fixed Point

```
CREATE TABLE teste(  
    x NUMERIC(1,0)    – same as NUMERIC(1)  
);
```

```
INSERT INTO teste VALUES (0);
```

```
INSERT INTO teste VALUES (-1);
```

```
INSERT INTO teste VALUES (9);
```

```
INSERT INTO teste VALUES (9.1);
```

⚠ rounded

```
INSERT INTO teste VALUES (11);
```

✗ error

⚠ Inserting in a NUMERIC with a precision or scale larger than specified will result in an error or in a warning; the value is often truncated.

Behaviour of Fixed Point

```
CREATE TABLE teste(  
    x NUMERIC(4,2)  
);
```

```
INSERT INTO teste VALUES (0.0);
```

```
INSERT INTO teste VALUES (-1.2);
```

```
INSERT INTO teste VALUES (12.34);
```

```
INSERT INTO teste VALUES (12.345); ⚠ Rounds to 12.35
```


Behaviour of Float

```
CREATE TABLE test(  
  x float,  
  y float  
);
```

```
INSERT INTO test VALUES(1.2, 1.2);
```

```
SELECT x+y FROM test;
```

```
2.40000000953674316
```

```
SELECT (1.0/3.0)*3.0;
```

Date Types

DATE

Stores dates
with a resolution
of
1 day

4 bytes

4713 BC
to
5874897 AD

TIME

Stores a time
with a resolution
of
1 microsecond

8 bytes

00:00:00
to
24:00:00

TIMESTAMP

Stores instant
timestamps with
resolution of
1 microsecond

8 bytes

4713 BC
to
294276 AD

Intervals

- ▶ **interval** (16 bytes) – enables capturing the difference between dates, times, or timestamps

⚠ Dates and times cannot be added. However, we can add intervals to dates and intervals to hours.

Indicative Field Sizes

Person/Organisation

| Field | Database Type | Max Size | Min Size | Validation |
|------------------------------------|---------------|----------|----------|--------------------------|
| PERSON/ORGANIZATION DETAILS | | | | |
| Person Full Name | VARCHAR | 80 | | |
| Company Name | VARCHAR | 200 | | |
| Street Address | VARCHAR | 255 | | |
| City | VARCHAR | 30 | | |
| Postal Code | VARCHAR | 12 | 2 | |
| Phone Number | VARCHAR | 15 | 3 | ITU E.16 |
| Phone Extension | VARCHAR | 11 | | ITU E.16 |
| Language | CHAR | 3 | | ISO 639 |
| Country Name | VARCHAR | 70 | | ISO 3166-1 |
| Latitude | NUMERIC | 9,6 | | |
| Longitude | NUMERIC | 8,6 | | |

Finance

| Field | Database Type | Max Size | Min Size | Validation |
|--------------------|---------------|----------|----------|------------|
| FINANCE | | | | |
| VAT ID | VARCHAR | 20 | 1 | |
| IBAN | VARCHAR | 30 | | |
| Credit Card Number | NUMERIC | 16 | | |
| Money | NUMERIC | 16,4 | | |

Electronic Commerce

| Field | Database Type | Max Size | Min Size | Validation |
|----------------------|---------------|----------|----------|---|
| ELECTRONIC | | | | |
| E-mail Address | VARCHAR | 254 | 6 | IETF RFC 3696 Checking email addresses |
| Domain Name | VARCHAR | 253 | 4 | |
| URL | VARCHAR | 2083 | 11 | |
| IP address (incl V6) | VARCHAR | 45 | 11 | |
| GUID | char | 36 | | |

Social Networks

| Field | Database Type | Max Size | Min Size |
|--------------------------|---------------|----------|----------|
| SOCIAL NETWORK | | | |
| Facebook max name length | VARCHAR | 50 | |
| Youtube channel | VARCHAR | 20 | |
| Twitter max name length | VARCHAR | 15 | |

NULL VALUES

NULL is a special value that means, simultaneously,
unfilled / unknown / not applicable

Suppose that Daniel did not specify his birthdate,
we could write the INSERT statement:

```
INSERT INTO employee VALUES(11, 'Daniel', null);
```

and then query the database for his record:

```
SELECT * FROM employee  
WHERE eid=11;
```

| eid | name | bdate |
|-----|--------|-------|
| 11 | Daniel | |

⚠ The use of NULL is ambiguous and should be avoided.

NOT NULL Constraint

- In SQL all columns (not part of the key) by default may have null VALUES.
- To prevent columns from taking null VALUES, we must add the **NOT NULL** constraint in front of the data type:


<field> *<type>* **NOT NULL**

```
CREATE TABLE employee(  
    ssn NUMERIC(11),  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    PRIMARY KEY(ssn)  
);
```

- Most columns (most often all columns) should be **NOT NULL**
- Any columns that participate in the **PRIMARY KEY** are already **NOT NULL** (because they null is never a valid value on a key)

Translating Associations

M:N Associations




```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY(eid, did),  
    FOREIGN KEY (eid) REFERENCES employee(eid),  
    FOREIGN KEY (did) REFERENCES department(did)  
);
```

Captures any **valid** combination of **⟨eid, did⟩**

M:N Auto Association

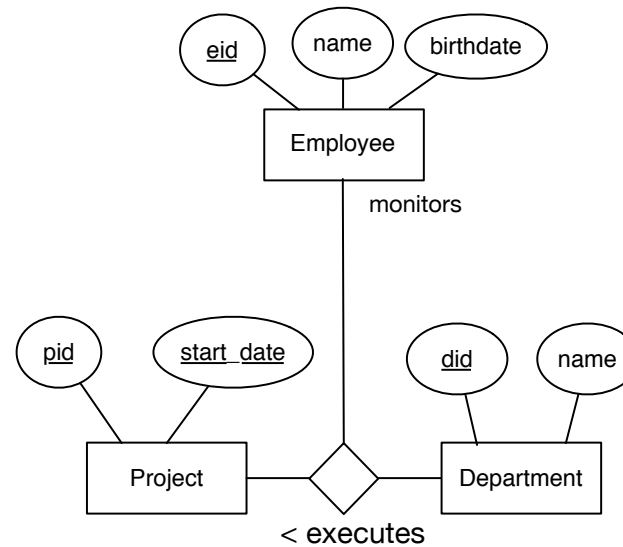
Special case of the self-association



```
CREATE TABLE supervises (  
    sup_eid INTEGER,  
    sub_eid INTEGER,  
    PRIMARY KEY(sup_eid, sub_eid)  
    FOREIGN KEY(sup_eid) REFERENCES  
        Employee(eid),  
    FOREIGN KEY(sub_eid) REFERENCES  
        Employee(eid)  
);
```

- Field names cannot repeat
- The fields *sup_eid* and *sub_eid* capture any valid combination of $\langle \text{eid}, \text{eid} \rangle$


Ternary Associations



```
CREATE TABLE executes (  
    eid INTEGER,  
    pid INTEGER,  
    did INTEGER,  
    PRIMARY KEY(eid, pid, did)  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(pid) REFERENCES project(pid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

Captures any *valid* combination of $\langle \text{eid}, \text{pid}, \text{did} \rangle$

One-to-many Associations




Guarantees that VALUES of *did* can *never repeat*!

```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY(did),  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

- Once a department is associated to an employee, *it cannot be associated again* (to another employee)
- We encode this by *guaranteeing* that each *did* appears *only once* in the table that represents the association (i.e., that *did* it is associated only once)

One-to-one Associations




Neither *did* nor *did* repeat
Therefore: Once a
'department' (or an 'employee')
exists in the table '*manages*', no
other can exist

```
CREATE TABLE Manages (  
  eid INTEGER,  
  did INTEGER,  
  since DATE,  
  PRIMARY KEY(eid),  
  UNIQUE (did),  
  NOT NULL(did),  
  FOREIGN KEY(eid) REFERENCES employee(eid),  
  FOREIGN KEY(did) REFERENCES department(did)  
);
```

- Both a Department and a Employee **can only** be associated **once**
- We encode this by **guaranteeing** that both **eid** and **did** appear **only once**

M:N Mandatory Participation




Every *department (did)* must participates in the '*manages*' association

```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since date,  
    PRIMARY KEY(eid, did),  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

⚠ There is no simple DBMS implementation for this constraint. Must often ensured by the application code.

Many-to-one Mandatory Participation



```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    bdate DATE NOT NULL,  
    since DATE NOT NULL,  
    did INTEGER NOT NULL,  
    PRIMARY KEY(eid)  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

- Instead of creating a table for the '*manages*' association, extend the table *employee* with the reference (a foreign key) to the department.
- Every record on *employee* must be *connected to* (in this case managing some) *department*