



# TÉCNICO LISBOA

## Sistemas de Informação e Bases de Dados 2020/2021

**Class 07: Translating E-A to SQL**

Prof. Paulo Carreira






# Class Outline

- ☐ Motivation
- ☐ Revisiting Referential Integrity
- ☐ Translating Entities and Attributes
- ☐ SQL Data Types
- ☐ Translating Column and Domain Constraints

# Motivation

# Translation Example



Employee		
eid	name	birthdate
1	Alice	10/10/1995
2	Bob	03/02/1996
3	Caroline	04/04/1997
4	Daniel	03/04/1998
5	Eduard	10/03/1994

works_at	
eid	did
1	1
2	2
3	3
4	3
5	3

Department		
did	name	location
1	Finance	Buraca
2	Marketing	Damaia
3	Sales	Chelas

⚠ **works\_at** encodes the association of Employees to Departments  
All VALUES must be coherent

# Relations and Attributes

```
CREATE TABLE employee(  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    bdate DATE NOT NULL,  
    PRIMARY KEY (eid)  
);
```

```
CREATE TABLE department(  
    did INTEGER,  
    name VARCHAR(20) NOT NULL,  
    location VARCHAR(20) NOT NULL,  
    PRIMARY KEY (did)  
);
```

```
INSERT INTO employee VALUES(1, 'Alice', '1995-10-10');  
INSERT INTO employee VALUES(2, 'Bob', '1996-03-02');
```

```
INSERT INTO department VALUES(1, 'Finance', 'Buraca');  
INSERT INTO department VALUES(2, 'Marketing', 'Damaia');
```

# Relations and Attributes

```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    bdate DATE NOT NULL,  
    PRIMARY KEY(eid)  
);
```

```
CREATE TABLE department (  
    did INTEGER,  
    name VARCHAR(20) NOT NULL,  
    location VARCHAR(20) NOT NULL,  
    PRIMARY KEY(did)  
);
```

```
CREATE TABLE works_at(  
    eid INTEGER,  
    did INTEGER,  
    PRIMARY KEY (eid, did)  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

Prevents inputting invalid *eid* or *did* values

```
INSERT INTO works_at VALUES(1, 1);  
INSERT INTO works_at VALUES(2, 1);  
INSERT INTO works_at VALUES(2, 99);
```

# Referential Integrity



# Table Relationships

A foreign key connects two tables

Foreign Key

Employee



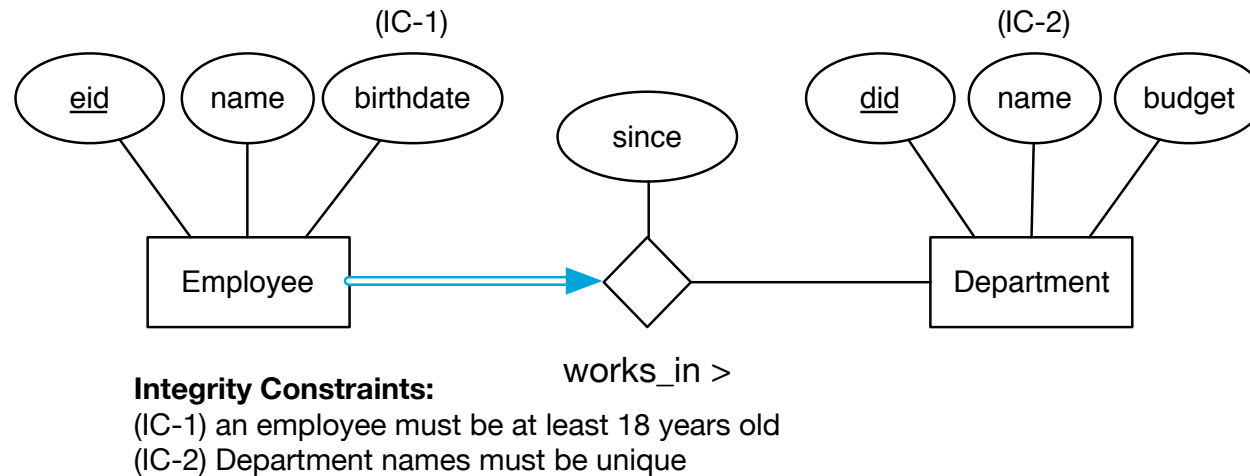
ID	Name	Tax id	T-Shirt	DID
001	João Guilherme Silva da Cunha	12345678	M	EN
002	Tomás Pinto dos Santos	91234567	M	EN
003	David Miguel Redwanz Duque	89012345	L	MK
004	Pedro Daniel Diz Pinela	67890123	M	HR
005	Guilherme de Queiróz Rebelo Brum Gomes	22394856	XL	EN
006	Marta Isabel de Almeida Cardoso	34562732	S	HR
007	Filipe Emanuel Lourenço Ramalho Fernandes	82533235	L	EN
008	Gabriel Filipe Queirós Mesquita Delgado Freire	23134539	M	EN
009	João Gomes Vultos Freitas	22231233	L	EN
010	Ricardo Afonso Rodrigues da Silva Oliveira	56372848	L	MK

Department

DID	Name	Budget
HR	Human Resources	50 000
EN	Software Engineering	1 200 000
MK	Marketing	150 000

Primary Key

# How to translate this case?



⚠ All VALUES in **Employee.dep** must exist as PRIMARY KEY the **Department.did**

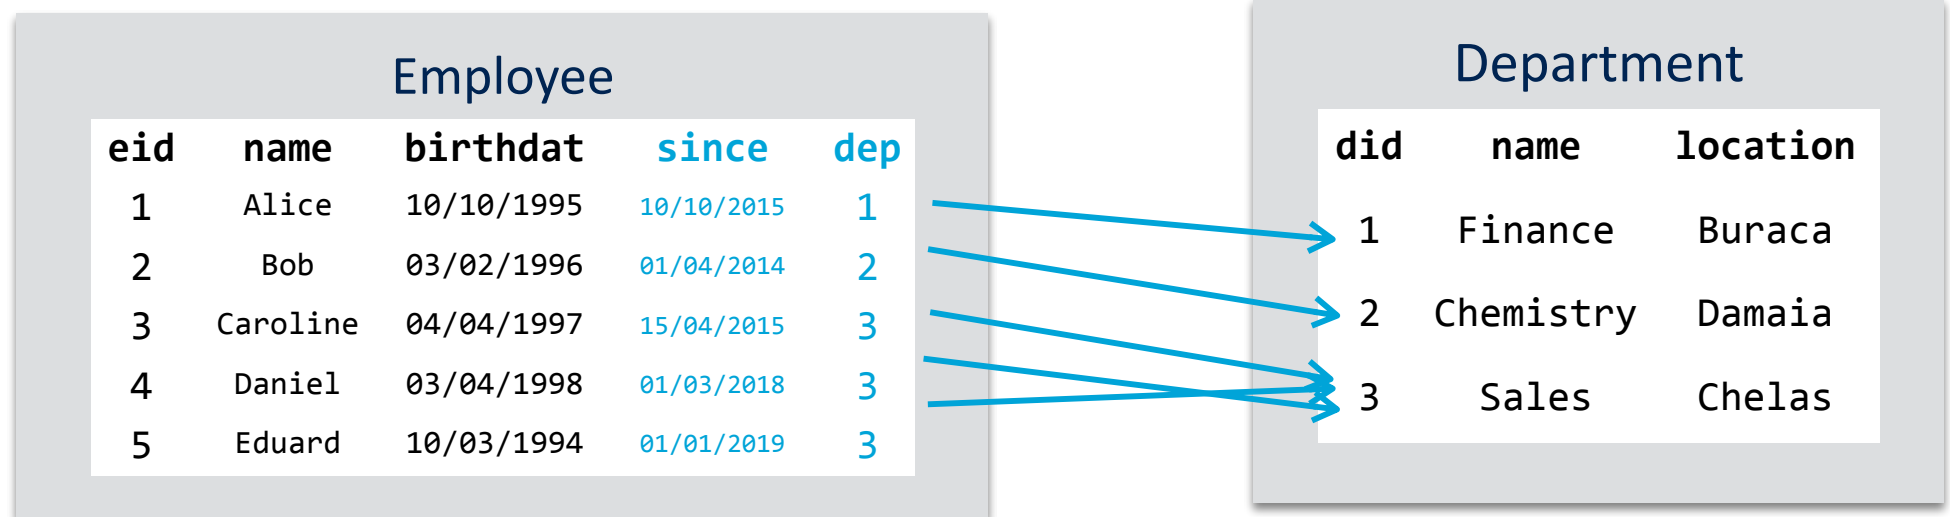
Employee

eid	name	birthdat	since	dep
1	Alice	10/10/1995	10/10/2015	1
2	Bob	03/02/1996	01/04/2014	2
3	Caroline	04/04/1997	15/04/2015	3
4	Daniel	03/04/1998	01/03/2018	3
5	Eduard	10/03/1994	01/01/2019	3

Department

did	name	location
1	Finance	Buraca
2	Chemistry	Damaia
3	Sales	Chelas

# Operations that violate Referential Integrity




- **Removing** lines from *Department*: We cannot remove *departments* to which *employees* are still associated
- **Updating VALUES** on *Department*: We cannot change VALUES on *department* that imply changing VALUES on *employee*
- **Inserting** lines on *Employee*: We cannot add *employees* on *departments* that do not exist

# Referential Integrity Constraints (or Foreign Keys)

- The most common constraint involving **two tables** is the **Referential Integrity** constraint
- Data in one table **must be always coherent with the data of** another table. A table is usually related to other tables.

# Translating Entities and Attributes

# Entities



```
CREATE TABLE employee (  
    eid    INTEGER,  
    name   VARCHAR(80) NOT NULL,  
    bdate  DATE NOT NULL,  
    PRIMARY KEY(eid)  
);
```

1. **Entities** result in a **table** with corresponding attributes
2. **PRIMARY KEY** is the same
3. **Constraints** have also to be translated (will see how later on...)

# Data Types and Values

# Basic Datatype Families

## Text Types

**Varchar**

**Char**

**Text**

'John Smith', 'R2D2', 'Red', ''

## Numeric Types

**Integer**

**Fixed Point**

**Floating Point**

-1, 25, +6.34, 0.5, 25e-03

## Date & Time

**Date**

**Time**

**Timestamp**

'2029-01-01', '08-JAN-2029 10:35:02'



# Text Types

## **VARCHAR(*n*)**

Variable length  
character string  
max size *n*

$n < 4000$

## **CHAR(*n*)**

Character string  
with fixed size *n*

$n < 4000$

## **TEXT**

Variable length  
text (multi-line)  
field (typically  
limited to 65535  
characters)

length typically  
limited to **65535**  
characters

# Integer Numeric Types

## INTEGER

An integer with  
up to 9 digits

4 bytes

$-2^{31}$   
-2 147 483 648  
to  
 $+2^{31}-1$   
2 147 483 647

## SMALLINT

An integer with  
up to 4 digits

2 bytes

$-2^{15}$   
-32 768  
to  
 $+2^{15}-1$   
32 767

## BIGINT

An integer with  
up to 18 digits

8 bytes

$-2^{63}$   
-9 223 372 036  
854 775 808  
to  
 $+2^{63}-1$   
9 223 372 036  
854 775 807

# Fixed Point Number Types

## **NUMERIC( $p, s$ )**

A numeric value  
with arbitrary  
exact precision

$\approx (3 + p/4)$  bytes

Up to **131072** digits  
before the decimal  
point

Up to **16383** digits  
after the decimal  
point

- Precision  **$p$** : total number of digits, must be positive ( **$p$**  > 0)
- Scale  **$s$** : number of digits to right of the decimal point, can be zero but must always be smaller than the precision ( $0 < \mathbf{s} < \mathbf{p}$ )

Whenever  $s$  is zero, we can  
write **NUMERIC( $p$ )**

# Floating Point Numeric Types

## REAL

variable-  
precision,  
inexact  
6 digits

4 bytes

$1\text{E}-37$   
to  
 $1\text{E}+37$

## DOUBLE

variable-  
precision,  
inexact  
15 digits

8 bytes

$1\text{E}-307$   
to  
 $1\text{E}+308$

# Fixed- vs. Floating Point

	FIXED POINT	FLOATING POINT
Precision	Fixed precision	Variable-precision
Storage	Stores numbers exactly	Stores numbers inexactly (as approximations)
Retrieval	Retrieved values never show any discrepancies	Retrieved values may show discrepancies from values stored
Speed	Slower calculations (*)	Faster calculations
Money Amounts	Can be safely used for monetary values	Should never be used for monetary values

(\*) In practice, this difference is often neglectable

# Behaviour of Fixed and Float

# Behaviour of Fixed Point

```
CREATE TABLE teste(  
    x NUMERIC(1,0)    -- same as NUMERIC(1)  
);
```

```
INSERT INTO teste VALUES (0);
```

```
INSERT INTO teste VALUES (-1);
```

```
INSERT INTO teste VALUES (9);
```

```
INSERT INTO teste VALUES (9.1);
```

 rounded

```
INSERT INTO teste VALUES (11);
```

 error

 Inserting in a NUMERIC with a precision or scale larger than specified will result in an error or in a warning; the value is often truncated.

# Behaviour of Fixed Point

```
CREATE TABLE teste(  
    x NUMERIC(4,2)  
);
```

```
INSERT INTO teste VALUES (0.0);
```

```
INSERT INTO teste VALUES (-1.2);
```

```
INSERT INTO teste VALUES (12.34);
```

```
INSERT INTO teste VALUES (12.345); ⚠ Rounds to 12.35
```



# Behaviour of Float

```
CREATE TABLE test(  
  x float,  
  y float  
);
```

```
INSERT INTO test VALUES(1.2, 1.2);
```

```
SELECT x+y FROM test;
```

```
2.40000000953674316
```

```
SELECT (1.0/3.0)*3.0;
```

# Date Types

## DATE

Stores dates  
with a resolution  
of  
1 day

4 bytes

4713 BC  
to  
5874897 AD

## TIME

Stores a time  
with a resolution  
of  
1 microsecond

8 bytes

00:00:00  
to  
24:00:00

## TIMESTAMP

Stores instant  
timestamps with  
resolution of  
1 microsecond

8 bytes

4713 BC  
to  
294276 AD

# Intervals

## INTERVAL

Enables capturing the difference between dates, times, or timestamps

4 bytes

-178 000 000 yrs  
to  
178 000 000 yrs

- Dates and times cannot be added (only subtracted)
- Intervals can be added to a date or to a time

DATE - DATE = INTERVAL

TIME - TIME = INTERVAL

DATE + INTERVAL = DATE

TIME + INTERVAL = TIME

# Behaviour of Interval

```
CREATE TABLE test(  
  a DATE,  
  b TIME,  
  c TIMESTAMP  
);
```

```
INSERT INTO test VALUES ('2020-12-12', '09:00', '2020-12-12 10:30');
```

```
SELECT * FROM test;
```

a		b		c
2020-12-12		09:00:00		2020-12-12 10:30:00

```
SELECT a + INTERVAL '10 days' AS a ,  
       b + INTERVAL '1 hour 30 min' AS b,  
       c + INTERVAL '20 days 2 hour 30 min' AS c  
FROM test;
```

a		b		c
2020-12-22 00:00:00		10:30:00		2021-01-01 13:00:00

# Indicative Field Sizes

# Person/Organisation

Field	Database Type	Max Size	Min Size	Validation
<b>PERSON/ORGANIZATION DETAILS</b>				
Person Full Name	VARCHAR	80		
Company Name	VARCHAR	200		
Street Address	VARCHAR	255		
City	VARCHAR	30		
Postal Code	VARCHAR	12	2	
Phone Number	VARCHAR	15	3	<a href="#">ITU E.16</a>
Phone Extension	VARCHAR	11		<a href="#">ITU E.16</a>
Language	CHAR	3		ISO 639
Country Name	VARCHAR	70		ISO 3166-1
Latitude	NUMERIC	9,6		
Longitude	NUMERIC	8,6		

# Finance

Field	Database Type	Max Size	Min Size	Validation
FINANCE				
VAT ID	VARCHAR	20	1	
IBAN	VARCHAR	30		
Credit Card Number	NUMERIC	16		
Money	NUMERIC	16,4		

# Electronic Commerce

Field	Database Type	Max Size	Min Size	Validation
ELECTRONIC				
E-mail Address	VARCHAR	254	6	IETF RFC 3696 Checking email addresses
Domain Name	VARCHAR	253	4	
URL	VARCHAR	2083	11	
IP address (incl V6)	VARCHAR	45	11	
GUID	char	36		



# Social Networks

Field	Database Type	Max Size	Min Size
SOCIAL NETWORK			
Facebook max name length	VARCHAR	50	
Youtube channel	VARCHAR	20	
Twitter max name length	VARCHAR	15	

# NULL Values

# NULL VALUES

**NULL** is a special value that means, simultaneously,  
unfilled / unknown / not applicable

Suppose that Daniel did not specify his birthdate,  
we could write the INSERT statement:

```
INSERT INTO employee VALUES(11, 'Daniel', null);
```

and then query the database for his record:

```
SELECT * FROM employee  
WHERE eid=11;
```

eid	name	bdate
11	Daniel	

⚠ The use of NULL is ambiguous and should be avoided.

# NOT NULL Constraint

- In SQL all columns (not part of the key) by default may have null VALUES.
- To prevent columns from taking null VALUES, we must add the **NOT NULL** constraint in front of the data type:

*<field>* *<type>* **NOT NULL**

```
CREATE TABLE employee(  
    ssn NUMERIC(11),  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    PRIMARY KEY(ssn)  
);
```


- Most columns (most often all columns) should be **NOT NULL**
- Any columns that participate in the **PRIMARY KEY** are already **NOT NULL** (because they null is never a valid value on a key)

# Translating Column and Domain Constraints

# Column Constraints


# PRIMARY KEY Constraints

## One Attribute



```
CREATE TABLE <table_name>(  
  a INTEGER,  
  b VARCHAR(80),  
  c NUMERIC(12,4),  
  d DATE,  
  PRIMARY KEY(a)  
);
```

## Multiple attributes




```
CREATE TABLE <table_name>(  
  a INTEGER,  
  b VARCHAR(80),  
  c NUMERIC(12,4),  
  d DATE,  
  PRIMARY KEY(a, b)  
);
```

A PRIMARY KEY constraint is specified as

**PRIMARY KEY**(*col*<sub>1</sub>, ..., *col*<sub>*n*</sub>)

# Uniqueness Constraints


## One "unique" attribute



**Integrity Constraints:**  
(IC-1) b is unique

```
CREATE TABLE table_name(  
    a INTEGER,  
    b VARCHAR(80),  
    c NUMERIC(12,4),  
    d DATE,  
    PRIMARY KEY(a),  
    UNIQUE(b)  
);
```

## Combination of "unique"



**Integrity Constraints:**  
(IC-1) (b,c) is unique  
(IC-2) (c,d) is unique

```
CREATE TABLE table_name(  
    a INTEGER,  
    b VARCHAR(80),  
    c NUMERIC(12,4),  
    d DATE,  
    PRIMARY KEY(a),  
    UNIQUE(b, c),  
    UNIQUE(c, d)  
);
```

A uniqueness constraint is specified as:

**UNIQUE**(*col*<sub>1</sub>, ..., *col*<sub>*n*</sub>)



# Domain Constraints

# Domain Constraint Checking

The **CHECK** clause can be used to specify the verification of the VALUES of any field of a record every time the record is inserted or updated:

**CHECK**(*condition*)

```
CREATE TABLE products (  
    product_no INTEGER,  
    name VARCHAR(80),  
    price NUMERIC,  
    discounted_price NUMERIC,  
    CHECK (price > 0),  
    CHECK (discounted_price > 0),  
    CHECK (price > discounted_price)  
);
```

# Domain Constraints

A **Domain constraint** guarantees that the VALUES of a column (field VALUES) are within the intended domain

```
CREATE TABLE employee(  
    ssn NUMERIC(11),  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    gender CHAR(1),  
    PRIMARY KEY(ssn),  
    CHECK (length(name) > 3),  
    CHECK (birthdate > '1920-01-01'),  
    CHECK (gender in ('M', 'F'))  
);
```

# Domain Constraint validation using a Technical Table

Whenever the domain is too large, the valid VALUES can be validated against a technical table

```
CREATE TABLE employee
(
  ssn NUMERIC(11),
  name VARCHAR(80) NOT NULL,
  birthdate DATE NOT NULL,
  birth_country CHAR(80),
  PRIMARY KEY(ssn),
  CHECK(birth_country
        IN (SELECT name FROM country))
);
```

Can be modelled as FOREIGN KEY

# Record/Line Constraints

A **record (row or line) constraint** is one that guarantees that the data of the record (row or line) is correct coherent

```
CREATE TABLE employee(  
    eid NUMERIC(9),  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    graduation DATE NOT NULL,  
    PRIMARY KEY(eid),  
    CHECK (LENGTH(name) > 3),  
    CHECK (birthdate > '1920-01-01'),  
    CHECK (extract(year FROM age(birthdate)) > 18),  
    CHECK graduation > birthdate  
);
```



# TÉCNICO LISBOA

## Sistemas de Informação e Bases de Dados 2020/2021

**Class 08: Translating E-A to SQL (cont)**

Prof. Paulo Carreira






# Class Outline

- ☐ Translating Associations
- ☐ Translating Specialisation/Generalisation
- ☐ Translating Weak Entities
- ☐ Translating Aggregations



# Translating Associations

# M:N Associations

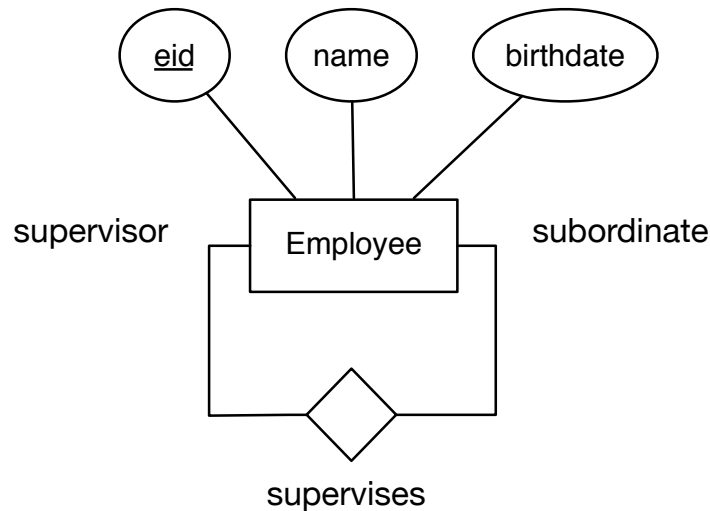


```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY(eid, did),  
    FOREIGN KEY (eid) REFERENCES employee(eid),  
    FOREIGN KEY (did) REFERENCES department(did)  
);
```

Captures any **valid** combination of **⟨eid, did⟩**

# M:N Auto Association


Special case of the self-association



```
CREATE TABLE supervises (  
    sup_eid INTEGER,  
    sub_eid INTEGER,  
    PRIMARY KEY(sup_eid, sub_eid)  
    FOREIGN KEY(sup_eid) REFERENCES  
        Employee(eid),  
    FOREIGN KEY(sub_eid) REFERENCES  
        Employee(eid)  
);
```

- Field names cannot repeat
- The fields *sup\_eid* and *sub\_eid* capture any valid combination of  $\langle \text{eid}, \text{eid} \rangle$

# Ternary Associations




Same translation as the case for the binary but with more legs

```
CREATE TABLE executes (  
    eid INTEGER,  
    pid INTEGER,  
    did INTEGER,  
    PRIMARY KEY(eid, pid, did)  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(pid) REFERENCES project(pid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

Captures any *valid* combination of  $\langle \text{eid}, \text{pid}, \text{did} \rangle$

# M:1 Associations




Guarantees that VALUES of *did* can never repeat!

```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY(did),  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

- Once a department is associated to an employee, it cannot be associated again (to another employee)
- We encode this by guaranteeing that each *did* appears only once in the table that represents the association (i.e., that *did* it is associated only once )

# 1:1 Associations




Neither *did* nor *did* repeat  
Therefore: Once a  
'department' (or an 'employee')  
exists in the table '*manages*', no  
other can exist

```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY(eid)  
    UNIQUE (did),  
    NOT NULL(did),  
    FOREIGN KEY(eid) REFERENCES employee(eid),  
    FOREIGN KEY(did) REFERENCES department(did)  
);
```

- Both a Department and a Employee **can only** be associated **once**
- We encode this by **guaranteeing** that both **eid** and **did** appear **only once**

# M:N Mandatory Participation




Every *department (did)* must participate in the '*manages*' association

```
CREATE TABLE manages (  
    eid INTEGER,  
    did INTEGER,  
    since date,  
    PRIMARY KEY(eid, did),  
    FOREIGN KEY(eid) REFERENCES  
        employee(eid),  
    FOREIGN KEY(did) REFERENCES  
        department(did)  
);
```

```
CREATE TABLE department(  
    did INTEGER,  
    name VARCHAR(20) NOT NULL,  
    location VARCHAR(20) NOT NULL,  
    PRIMARY KEY (did)  
    -- Every department must exist  
    in the table 'manages'  
);
```

⚠ There is no simple DBMS implementation for this constraint. Must often be ensured by the application code.

# M:1 Mandatory Participation



```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    since DATE NOT NULL,  
    manages_did INTEGER NOT NULL,  
    PRIMARY KEY(eid)  
    FOREIGN KEY(manages_did) REFERENCES department(did)  
);
```

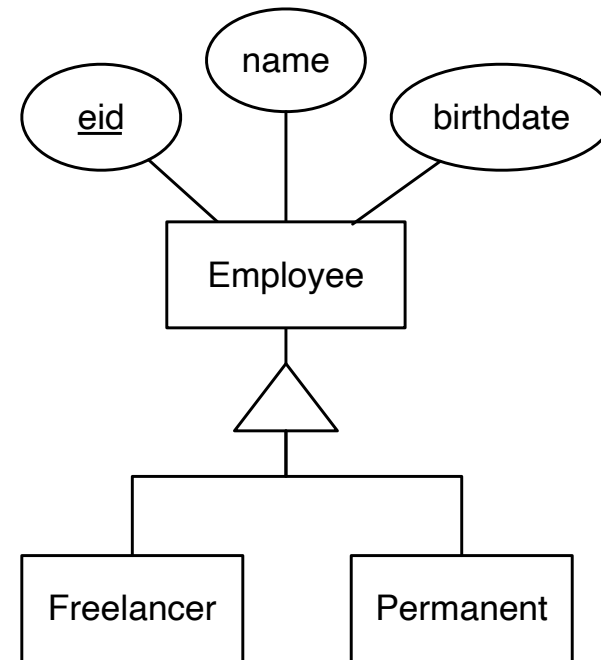
- Instead of creating a table for the '*manages*' association, we extend the table *employee* with the reference (a foreign key) to the department.
- Every record on *employee* must be *connected to* (in this case managing some) *department*



# Translating Generalisations

# Simple Specialisation

```
CREATE TABLE employee (  
  eid INTEGER,  
  name VARCHAR(80) NOT NULL,  
  birthdate DATE NOT NULL,  
  PRIMARY KEY(empid)  
);
```

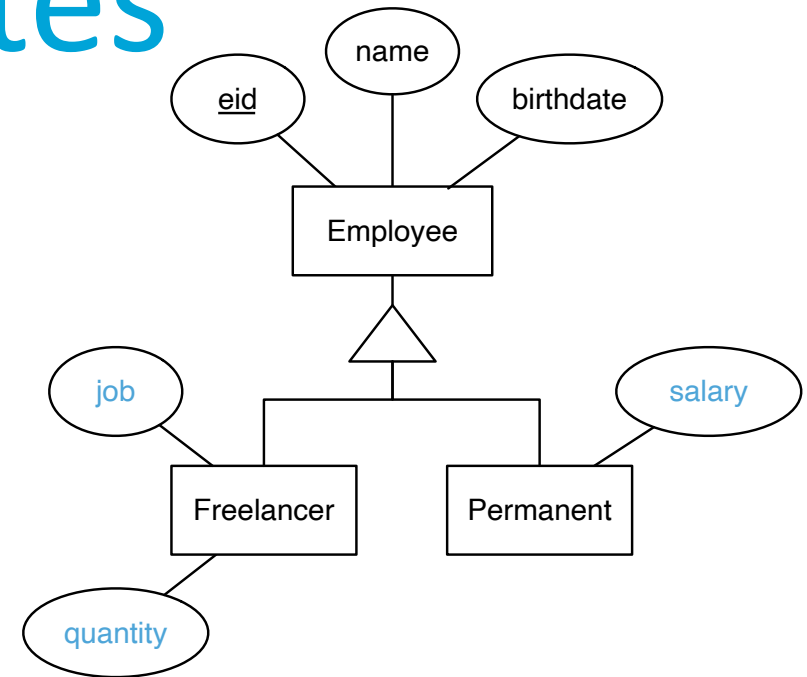


```
CREATE TABLE freelancer (  
  eid INTEGER,  
  PRIMARY KEY(eid),  
  FOREIGN KEY(eid) REFERENCES  
    employee(eid)  
);
```

```
CREATE TABLE contracted (  
  eid INTEGER,  
  PRIMARY KEY(eid),  
  FOREIGN KEY (eid) REFERENCES  
    employee(eid)  
);
```

# Specialisation with attributes

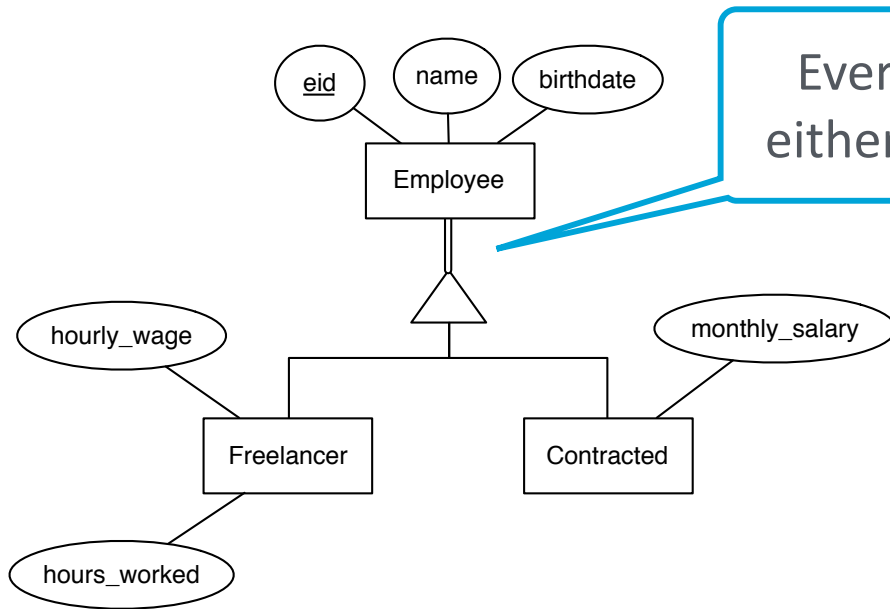
```
CREATE TABLE employee (  
  eid INTEGER,  
  name VARCHAR(80) NOT NULL,  
  birthdate DATE NOT NULL,  
  PRIMARY KEY(empid)  
);
```



```
CREATE TABLE freelancer (  
  eid INTEGER,  
  hourly_wage NUMERIC(12,4),  
  hours_worked INTEGER,  
  PRIMARY KEY(eid),  
  FOREIGN KEY(eid) REFERENCES  
    employee(eid)  
);
```

```
CREATE TABLE contracted (  
  eid INTEGER,  
  salary NUMERIC(12,4),  
  PRIMARY KEY(eid),  
  FOREIGN KEY (eid) REFERENCES  
    employee(eid)  
);
```

# Mandatory Specialisation



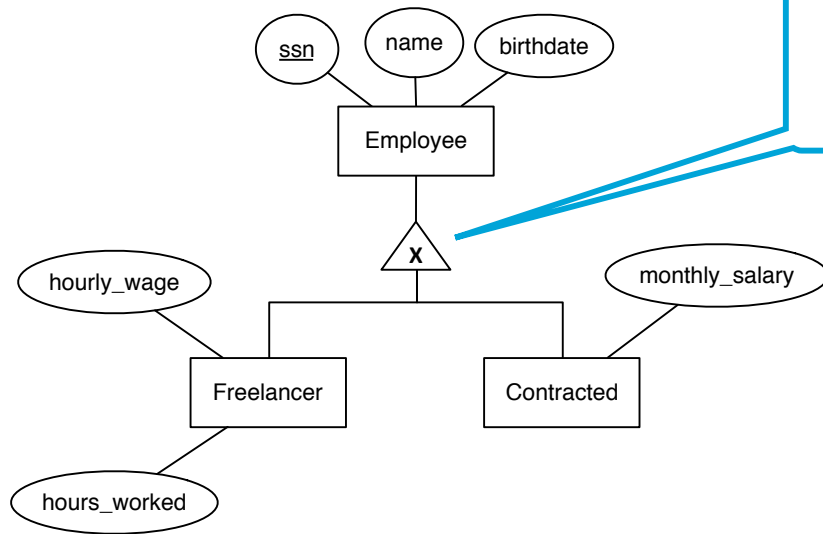
Every **eid** of **employee** must exist either in **freelancer** or in **contracted**

```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    PRIMARY KEY(empid)  
    -- Every employee must exist either  
    -- in the table 'freelancer' or in  
    -- the table 'permanent'  
);
```

```
CREATE TABLE freelancer (  
    eid INTEGER,  
    hourly_wage NUMERIC(12,4),  
    hours_worked INTEGER,  
    PRIMARY KEY(eid),  
    FOREIGN KEY(eid) REFERENCES  
        employee(eid)  
);
```

```
CREATE TABLE contracted (  
    eid INTEGER,  
    monthly_salary NUMERIC(12,4),  
    PRIMARY KEY(eid),  
    FOREIGN KEY (eid) REFERENCES  
        employee(eid)  
);
```

# Disjoint Specialisation



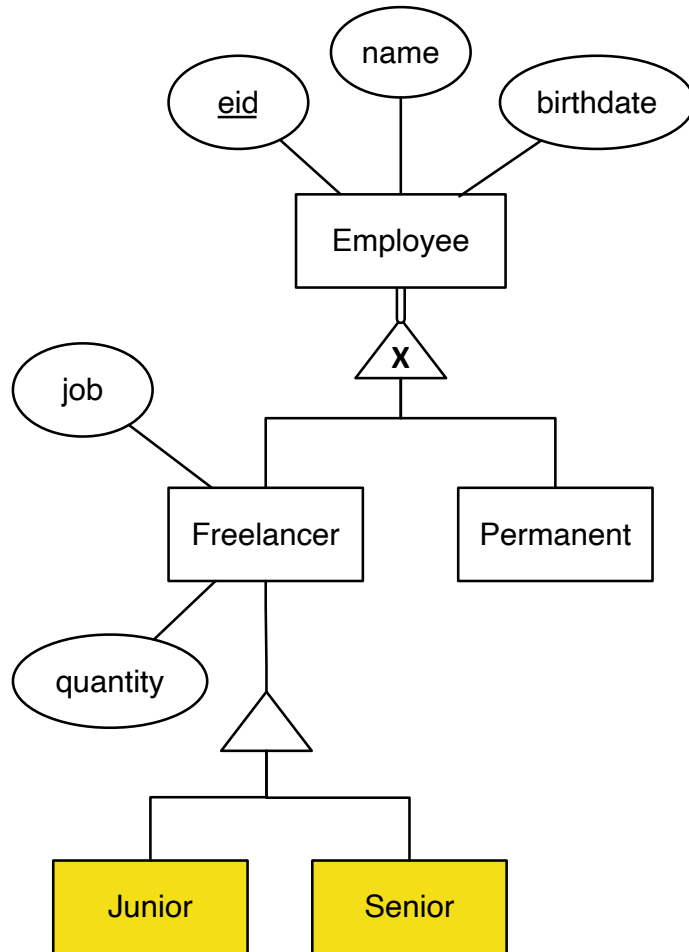
The **eid** of an Employee cannot exist in Freelancer and Contracted at the same time

```
CREATE TABLE employee (  
    eid INTEGER,  
    name VARCHAR(80) NOT NULL,  
    birthdate DATE NOT NULL,  
    PRIMARY KEY(eid)  
    -- No employee can exist at the same  
    -- time in the both the table  
    -- 'freelancer' or in the table  
    -- 'permanent'  
);
```

```
CREATE TABLE freelancer (  
    eid INTEGER,  
    hourly_wage NUMERIC(12,4),  
    hours_worked INTEGER,  
    PRIMARY KEY(eid),  
    FOREIGN KEY(eid) REFERENCES  
        employee(eid)  
);
```

```
CREATE TABLE contracted (  
    eid INTEGER,  
    monthly_salary NUMERIC(12,4),  
    PRIMARY KEY(eid),  
    FOREIGN KEY (eid) REFERENCES  
        employee(eid)  
);
```

# Nested Specialisation



```
create table junior (
    eid INTEGER,
    PRIMARY KEY (eid),
    FOREIGN KEY (eid) REFERENCES
        freelancer(eid)
);
```

```
CREATE TABLE employee (
    eid INTEGER,
    name VARCHAR(80) NOT NULL,
    bdate DATE NOT NULL,
    PRIMARY KEY(empid)
    -- <Mandatory constraint...>
    -- <Disjoint constraint...>
);
```

```
CREATE TABLE freelancer(
    eid INTEGER,
    hourly_wage money,
    hours_worked INTEGER,
    PRIMARY KEY(eid),
    FOREIGN KEY(eid) REFERENCES
        employee(eid)
);
```

```
CREATE TABLE contrated(
    eid INTEGER,
    monthly_salary money,
    PRIMARY KEY (eid),
```

```
CREATE TABLE senior (
    eid INTEGER,
    PRIMARY KEY (eid),
    FOREIGN KEY (eid) REFERENCES
        freelancer(eid)
);
```

# Mapping Generalisations/ Specialisations


1. Map the super-entity in a table
2. Map sub-entities into distinct tables where:
  - ▶ The **key** of each table corresponding to a **sub-entity** is the **key** of the **super-entity** (enforced with the corresponding FK constraint)
3. Disjoint or Mandatory specialisation constraints are identified as comments (they will mapped through ICs over the super-entity **using advanced database programming or application code**)

⚠ There is no simple DBMS implementation for **Disjoint** and **Mandatory** constraints. This will be explained later in the course.

# Translating Weak Entities



# Weak Entities



The translation is similar to the case of M:1 Mandatory Participation

```
CREATE TABLE person(  
  name VARCHAR(80),  
  ssn NUMERIC(9)  
  PRIMARY KEY(ssn)  
);
```

```
CREATE TABLE animal(  
  ssn NUMERIC(9)  
  name VARCHAR(80),  
  species VARCHAR(20),  
  breed VARCHAR(20),  
  PRIMARY KEY(ssn, name),  
  FOREIGN KEY(ssn) REFERENCES person(ssn)  
);
```

The PRIMARY KEY is the combination of the key of the strong entity with the partial key

# Weak Entities

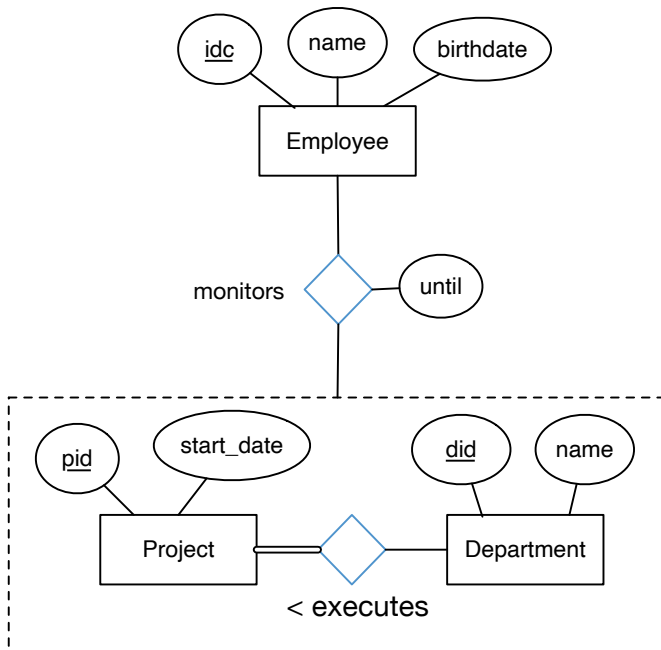
The **Weak Entities** originate a table that has a key composed by:

1. The association key that corresponds to the strong entity
2. The specified partial key
3. Any attributes of the weak entity (if they exist)

The association is not converted into a table

# Translating Aggregations

# Aggregation



```
CREATE TABLE executes(  
  pid INTEGER,  
  did INTEGER,  
  PRIMARY KEY (pid, did)  
  FOREIGN KEY (pid) REFERENCES project(pid)  
  FOREIGN KEY (did) REFERENCES department(did)  
);
```

```
CREATE TABLE monitors(  
  eid INTEGER,  
  pid INTEGER,  
  did INTEGER,  
  until date,  
  PRIMARY KEY (eid, pid, did),  
  FOREIGN KEY (pid, did) REFERENCES executes(pid, did)  
  FOREIGN KEY (eid) REFERENCES employee(eid)  
);
```

# Aggregation

An Aggregation is mapped as an association where:

1. The interior of the aggregation is mapped to a table
2. The association with the aggregation is mapped into a table