



# TÉCNICO LISBOA

## Sistemas de Informação e Bases de Dados


**Class 11: Triggers e Stored Procedures**

Prof. Paulo Carreira





# Bibliography



Chapters 3 e 5

## POSTGRES Manual

- <https://www.postgresql.org/docs/9.5/static/sql-createtrigger.html>
- <https://www.postgresql.org/docs/9.5/static/plpgsql-structure.html>
- <https://www.postgresql.org/docs/9.5/static/plpgsql-control-structures.html>

# Class Outline

- ☐ Introduction to PSM
- ☐ Stored Functions and Procedures
- ☐ Syntax of the main block elements
- ☐ Cursors

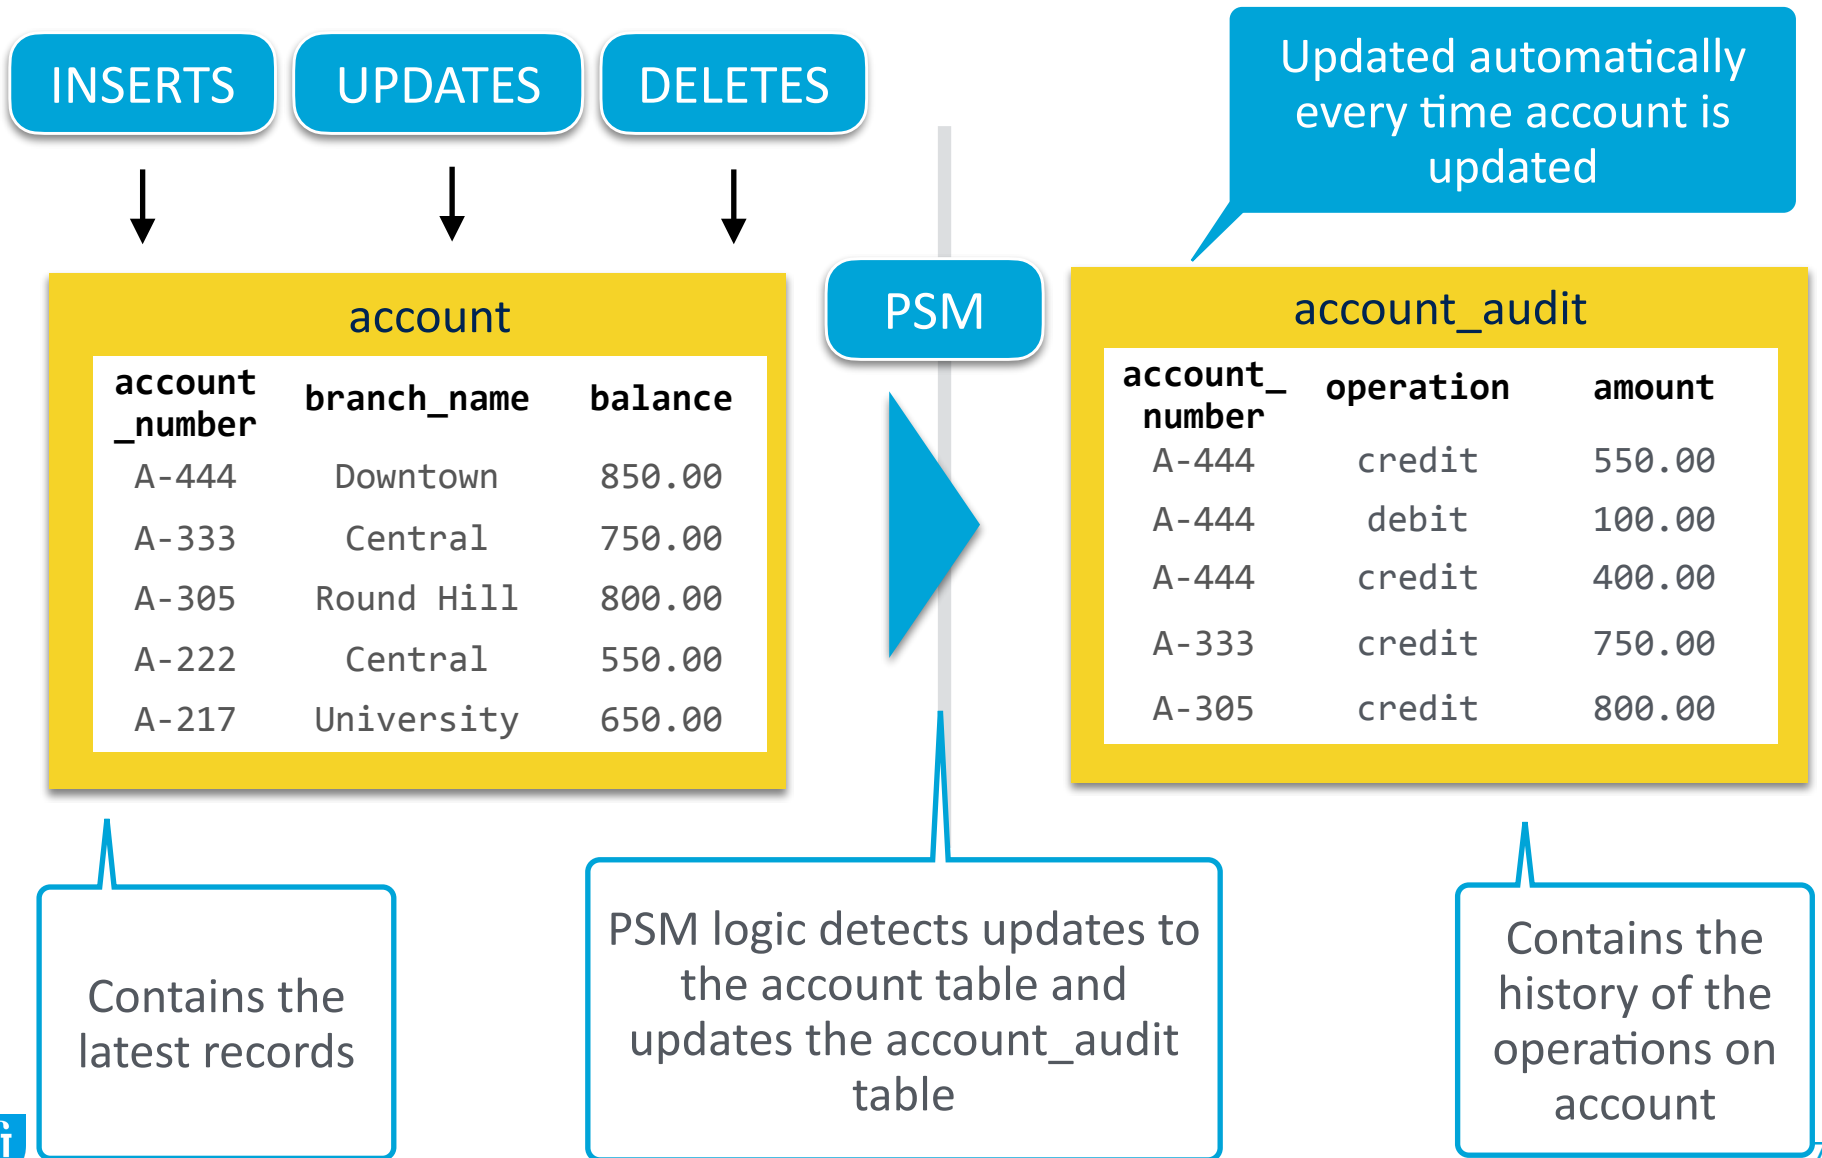
# Introduction

# Data Centric Development



**Data-centric database** development is a development approach whereby procedural instructions (like as a sub-program in any programming language) can be stored inside a DB also known as a **stored program** / **stored routine**

# Motivation Example



# Persistent Stored Modules




PSM

A Persistent Stored Module can be of two kinds:

- A *stored function* that returns a value but that should not change the state of the database.
- A *stored procedure* that may change the state of the database and does not return a value




# Stored Function



A **Stored Function** combines inputs and results of queries from the database to produce the output. The database is never changed.

# Stored Procedure



A **Stored Procedure** combines inputs and results of queries from the database to produce the output. The database is never changed.

# RDBMS procedural languages

- (c. 1992) **PL/SQL** – Oracle
- (c. 1998) **PL/pgSQL** – Postgres ( $\approx$ PL/SQL)
- (c. 1999) **SQL/PSM** – Standard SQL (IBM DB2 and MySQL)
- (c. 1989) **Transact-SQL** – MS SQL Server + Sybase

# Advantages

- Can **make applications faster**: PSM are compiled and maintained inside the database
- **Reduce the exchange of data**: especially between the application and the DB server
- **Introduce a level of indirection**: can be called from applications written in distinct languages

# Drawbacks

**Development Complexity:** Maintaining PSMs is relatively complex (especially when domain logic is complex)

- **Debugging** and **profiling** become very difficult
- Development environment highly dependent on the DBMS used and on the toolchain of the manufacturer

# Definition and invocation

- Invoking a procedure and a function

```
CREATE [OR REPLACE] PROCEDURE/FUNCTION ...
```

```
DROP [OR REPLACE] PROCEDURE/FUNCTION ...
```

- Defining a procedure or a function

```
CALL nome_procedimento [(param1, param2,...)]
```

```
my_function_name([param1, param2,...])
```

Stored functions and procedures **can be called** from other procedures, or from client applications written in C, PHP, Java, Javascript

# Statements (the body of a PSM)

# Example procedure (PSM Standard)

```
CREATE PROCEDURE my_procedure(IN parameter1 INTEGER) AS
```

```
DECLARE variable1 CHAR(10);
```

```
BEGIN
```

```
    IF parameter1 = 17
```

```
    THEN
```

```
        SET variable1 := 'birds';
```

```
    ELSE
```

```
        SET variable1 := 'beasts';
```

```
    END IF;
```

```
    INSERT INTO table1
```

```
        VALUES (variable1);
```

```
END
```

Name and parameters

Variable declarations

Start of the statement  
block

Assignment

SQL statement



# Stored Functions in POSTGRES

Language  
Token  
separator

```
CREATE FUNCTION myfunc([params])  
RETURNS type AS  
$$  
DECLARE [declarations]  
BEGIN  
    [statements]  
END  
$$ LANGUAGE plpgsql;
```

PL/PgSQL does not support true stored procedures;  
only **stored functions** are supported.

```
CREATE FUNCTION myfunc([params]) RETURNS VOID AS
```

PSM in Postgres does not follow the SQL PSM  
Standard (although they are very similar)

# The 'Hello World' of PSM

Define a simple function to add two numbers

```
CREATE FUNCTION add_me(  
    x NUMERIC,  
    y NUMERIC)  
    RETURNS NUMERIC AS  
$$  
BEGIN  
    RETURN x + y;  
END  
$$ LANGUAGE plpgsql;
```

# Executing a Stored Function

Distinct ways to execute a stored function

- Direct query (without FROM)

```
SELECT add_me(2,3);
```

```
add_me  
-----  
5
```

- As an expression in the SELECT (applied to every line)

```
SELECT name, add_me(salary, 100)  
FROM employee
```

- As an expression in the WHERE (applied to every line)

```
SELECT *  
FROM employee  
WHERE add_me(salary, -1000) < 5000
```

# Specifying Deterministic Functions

## Definition

A function is *deterministic* if it always produces the same results for the same input values

```
CREATE FUNCTION myfunc([params]) RETURNS type  
IMMUTABLE
```

- Example of a *non-deterministic* function: invokes the NOW( ) or RAND( ) function
- In POSTGRES, every functions are non-deterministic by default (we must specify **IMMUTABLE** for deterministic functions)
- Immutable (i.e. deterministic) functions can be optimised

# Variable Declarations

```
DECLARE var_name [, var_name ...] type  
      [DEFAULT value]
```

Variables are only visible in the scope of the **BEGIN ... END** block where they were declared

```
DECLARE  
    first_name VARCHAR(50) DEFAULT 'John';  
    last_name  VARCHAR(50) DEFAULT 'Doe';  
    counter    INTEGER := 1;  
    payment    NUMERIC(11,2) := 20.5;
```

# Example

*A function that, given the **name of a depositor**, returns how many accounts that depositor has*

```
CREATE OR REPLACE FUNCTION account_count(  
    d_name VARCHAR(40))  
RETURNS INTEGER AS  
$$  
DECLARE total_count INTEGER;  
BEGIN  
    SELECT COUNT(*) INTO total_count  
    FROM depositor  
    WHERE customer_name = d_name;  
  
    RETURN total_count;  
END  
$$ LANGUAGE plpgsql;
```

# Example Usage

*Get the name, street and city of customers with  
more than one account*

```
SELECT
    customer_name,
    customer_street,
    customer_city
FROM customer
WHERE account_count(customer_name) > 1
```

Demonstration example only.  
Do not use in practice!

# Assigning values to local variables

- Assigning values from expressions

```
var_name := expr
```

POSTGRES syntax

```
LET var_name := expr
```

```
SET var_name := expr
```

Other RDBMS

- Capturing the result of query

```
SELECT col_name, ...  
INTO var_name  
FROM ...
```



# Conditionals and Cycles

```
IF condition THEN statement_list1  
ELSE statement_list2  
END IF;
```

```
WHILE condition  
DO statement_list  
END WHILE;
```

```
REPEAT statement_list  
UNTIL condition  
END REPEAT;
```

```
LOOP statement_list  
END LOOP;
```

# Examples

# A PSM that returns a table

*A function that, given the name of the customer, returns all customer accounts*

```
CREATE FUNCTION accounts_of(name VARCHAR(80))  
  RETURNS SETOF account  
AS  
$$  
  SELECT a.account_number, branch_name, balance  
  FROM account as a, depositor as d  
  WHERE a.account_number = d.account_number  
        and d.customer_name = name;  
$$  
LANGUAGE sql;
```

# Function with complex types

*Function that, given the name of a customer, returns multiple rows (a table) with the customer's accounts*

```
DROP TYPE account_data;
```

```
CREATE TYPE account_data AS (  
    account_number CHAR(5),  
    branch_name    VARCHAR(80),  
    balance        NUMERIC(16,4)  
);
```

```
DROP FUNCTION accounts_of(varchar);
```

```
CREATE FUNCTION accounts_of(name VARCHAR(80))  
    RETURNS account_data  
AS $$  
    SELECT account_number, branch_name, balance  
    FROM account NATURAL JOIN depositor  
    WHERE customer_name = name;  
$$ LANGUAGE sql;
```

```
SELECT * FROM accounts_of('Cook');
```

```
SELECT accounts_of('Cook'); -- Works but gives different result
```

# A PSM that returns a row of a table

```
CREATE OR REPLACE FUNCTION accounts(  
    acct_num VARCHAR(10))  
RETURNS account AS  
$$  
DECLARE  
    acct account%ROWTYPE;  
BEGIN  
    SELECT *  
    INTO acct  
    FROM account  
    WHERE acct_number = acct_num;  
  
    RETURN acct;  
END  
$$  
LANGUAGE plpgsql;
```

# Syntax for "Stored Procedures"

```
CREATE FUNCTION myfunc([params]) RETURNS type
```

- Only have stored functions are supported in POSTGRES:  
The **RETURNS** statement must be specified
- Procedures are functions that do not return anything:  
Obtained by specifying **RETURNS VOID**

<https://www.postgresql.org/docs/9.2/static/sql-createfunction.html>

# DO blocks

**DO** executes an anonymous code block

```
DO $$  
DECLARE total NUMERIC DEFAULT 0;  
BEGIN  
    SELECT SUM(balance)  
    INTO total  
    FROM account;  
  
    RAISE INFO 'Account total --> %', total;  
END$$;
```

```
INFO: Account total --> 6300.0000
```

The code block is treated as though it were the body of a function with no parameters returning void

# Cursors

Abstraction to **read a table with a file semantics** that can be used within Persistent Stored Modules

- **OPEN** - Open the cursor (and execute the query associated to the cursor)
- **FETCH** - Read the record and advance to the next
- **CLOSE** - Close the cursor and free resources



⚠ Never use floating-point data types for money

# Cursors

```
CREATE OR REPLACE FUNCTION average_balance() RETURN FLOAT
AS
$$
  DECLARE balance_var REAL DEFAULT 0.0;
  DECLARE sum_balance REAL DEFAULT 0.0;
  DECLARE count_balance INTEGER DEFAULT 0;
  DECLARE cursor_account CURSOR FOR
    SELECT balance FROM account;
BEGIN
  OPEN cursor_account;
  LOOP
    FETCH cursor_account INTO balance_var;
    sum_balance := sum_balance + balance_var;
    count_balance := count_balance + 1;
  END LOOP;
  CLOSE cursor_account;
  RETURN sum_balance / count_balance;
END
$$ LANGUAGE plpgsql;
```

Declares the cursor for a query

Opens the cursor

Reads a record and advances to the next record

Closes the cursor



# TÉCNICO LISBOA

## Sistemas de Informação e Bases de Dados

**Class 12: Triggers, Stored Procedures, Views**

Prof. Paulo Carreira







slides não são livros

# Class Outline

- Exceptions
- Triggers
- Views

# Exceptions

# Raising Exceptions

**RAISE EXCEPTION** *<message>*  
**USING HINT** '*text*'

```
CREATE FUNCTION search(uid INT) RETURNS VOID AS
$$
BEGIN
    IF NOT EXISTS(
        SELECT *
        FROM users
        WHERE user_id=uid)
    THEN
        RAISE EXCEPTION 'Nonexistent ID %', user_id
        USING HINT = 'Please check your user ID';
    END IF;
END;
$$ LANGUAGE plpgsql;
```

# Recovering from Exceptions

```
CREATE OR REPLACE FUNCTION calc_tax(arg NUMERIC) RETURNS
NUMERIC
AS $$
DECLARE res INTEGER;
BEGIN
    -- main block
    res := 100 / arg;
    RETURN res;

    EXCEPTION
        WHEN division_by_zero
        THEN RETURN 0;
END;
$$
LANGUAGE plpgsql;
```

Some code that  
triggers the  
exception

Use the WHEN clause  
to catch the exception



# Triggers

# Triggers

Are **actions** (procedures) that are automatically invoked in response to specific database **events** (operations) such as inserts, updates, or deletes

- Implement complex Integrity Constraints
- Update/change tables (for history or audit purposes) when some other table is changed

# Triggers E-C-A pattern

Triggers are specified using *event-condition-action* (ECA) rules:

- **Event:** what type of update activates the trigger
- **Condition:** a question or test to see if the action should be taken
- **Action:** a procedure that is executed when the trigger is activated and the previous condition is true

# Syntax

A procedure that is automatically invoked in response to specific database updates

- Creating a Trigger

```
CREATE TRIGGER <trigger_name>  
    { BEFORE | AFTER } { INSERT | UPDATE | DELETE }
```

```
ON <tbl_name>
```

```
WHEN <condition>
```

```
FOR EACH { ROW | STATEMENT } EXECUTE PROCEDURE <proc_name>
```

- Removing a Trigger

```
DROP TRIGGER trigger_name
```

*Triggers are an old concept was only standardised in the SQL standard: 1999*

# Example: Integrity Constraint

*Create a trigger that prevents the balance from being negative or greater than 100*

```
CREATE OR REPLACE FUNCTION chk_balance_interval_proc()  
RETURNS TRIGGER AS  
$$  
BEGIN  
    IF NEW.balance < 0 OR NEW.balance > 100 THEN  
        RAISE EXCEPTION 'Withdrawal or deposit past the limits'  
    END IF;  
  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;
```

In POSTGRES  
we must  
create a  
function that  
returns a  
TRIGGER! 💡

This trigger  
may fail

```
CREATE TRIGGER chk_balance_interval  
BEFORE UPDATE OR INSERT ON account  
FOR EACH ROW EXECUTE PROCEDURE chk_balance_interval_proc();
```

```
UPDATE account  
SET balance = balance - 500  
WHERE account_number = 'A-101';
```

# Example: Integrity Constraint

*Create a trigger that never lets the balance go below 0 or above 100*

```
CREATE OR REPLACE FUNCTION chk_balance_interval_proc()  
RETURNS TRIGGER AS  
$$  
BEGIN  
    IF NEW.balance < 0 THEN  
        NEW.balance := 0;  
    ELSEIF NEW.balance > 100 THEN  
        NEW.balance := 100;  
    END IF;  
  
    RETURN NEW;  
END;  
$$ LANGUAGE plpgsql;  
  
-- Code to install the trigger should go here...
```

This trigger never fails;  
instead, it intercepts  
the updates to balance

All account  
values will  
remain  
between 0  
and 100

```
UPDATE account  
SET balance = balance - 500
```

# Trigger Failure Behaviour

- If a **BEFORE** trigger fails, the operation on the corresponding row or table is not performed
- An **AFTER** trigger is only executed if any **BEFORE** triggers on the same table and related to the same operation are successfully executed

# Example: Complex table update

*Consider the following requirements for the Bank database*

- ▶ Whenever a customer withdraws more than the balance
- Instead of resulting in a negative balance, the bank:
  1. Creates a loan equal to the missing amount
  2. Gives the loan the same number as the account
  3. Zeroes the account balance
- ▶ **Trigger condition:** an update that results in a negative account balance



# Example: Complex table update

*Withdrawal of € 800 from the A-102 account*

account_number	branch_name	balance
A-101	Downtown	500.0000
A-215	Metro	600.0000
A-102	Uptown	700.0000
A-305	Round Hill	800.0000
A-201	Uptown	900.0000
A-222	Central	550.0000
A-217	University	650.0000
A-333	Central	750.0000
A-444	Downtown	850.0000

# Trigger: Example 2

*Withdrawal of € 800 from the A-102 account result  
on a loan*

## Account

account_number	branch_name	balance
A-101	Downtown	500.0000
A-215	Metro	600.0000
A-102	Uptown	700.0000
A-305	Round Hill	800.0000
A-201	Uptown	900.0000
A-222	Central	550.0000
A-217	University	650.0000
A-333	Central	750.0000
A-444	Downtown	850.0000

## Loan

loan_number	branch_name	amount
L-17	Downtown	1000.0000
L-23	Central	2000.0000
L-15	Uptown	3000.0000
L-14	Downtown	4000.0000
L-93	Metro	5000.0000
L-11	Round Hill	6000.0000
L-16	Uptown	7000.0000
L-20	Downtown	8000.0000
L-21	Central	9000.0000
A-102	UpTown	100.0000



## Depositor

customer_name	account_number
Johnson	A-101
Brown	A-215
Cook	A-102
Cook	A-101
Flores	A-305
Johnson	A-201
Iacocca	A-217
Evans	A-222
Oliver	A-333
Brown	A-444

## Borrower

customer_name	loan_number
Iacocca	L-17
Brown	L-23
Cook	L-15
Nguyen	L-14
Davis	L-93
Brown	L-11
Gonzalez	L-17
Iacocca	L-16
Parker	L-20
Brown	L-21
Cook	A-102

# Trigger: Example 2

```
CREATE OR REPLACE FUNCTION overdraft_proc()  
RETURNS TRIGGER  
AS $$  
BEGIN  
    IF NEW.balance < 0 THEN  
        INSERT INTO loan  
        VALUES (NEW.account_number,  
                NEW.branch_name,  
                NEW.balance);  
  
        INSERT INTO borrower  
        SELECT customer_name, account_number  
        FROM depositor  
        WHERE depositor.account_number = NEW.account_number;  
  
        NEW.balance := 0;  
    END IF;  
END  
$$ LANGUAGE SQL  
  
CREATE TRIGGER overdraft_trigger  
BEFORE UPDATE ON account  
FOR EACH ROW EXECUTE PROCEDURE overdraft_proc();
```

# Problems with *Triggers*

- **Complex effect.** Its effect can be complex and unpredictable
  - *Multiple triggers* can be triggered in a single operation
  - The action of a trigger can activate another trigger (*recursive triggers*)
- **Cycles of events.** Chains of endless events. Very difficult to debug and fix.

# Problems with Triggers

- **Unintended executions**
  - Changes to a table will run triggers even if that's not what is intended
- **Occurrence of errors**
  - If the trigger fails, the entire operation fails
  - Extended undo recovery times

# When **not** to use Triggers

- Summary tables: Use *VIEWS* instead of triggers, if possible
- Complex Integrity Constraints: Use *CHECKs* whenever possible
- Table replication: Use DBMS build-in mechanisms

# Views

# Views

- A view is a *virtual table* defined through a query
- Associates a name to a **SELECT** statement

**CREATE VIEW** *myview* **AS SELECT ...**

Once created, it can be used as a *relation* but it is not the same as creating a table

**Views** do not have storage, they are computed, and their contents will change if the tables involved in the query change



# Creating Views

Views are database objects that can be *created* and *removed*.

- Creation of View

```
CREATE VIEW myview AS SELECT ...
```

- Removal of View

```
DROP VIEW myview
```

# Creating a View: Example 1

```
CREATE VIEW account_stats(name, num_accts)
AS
SELECT name, COUNT(*) AS num_accts
FROM depositor
GROUP BY customer_name;
```

```
SELECT * FROM account_stats;
```

# Resolving Queries on Views

The technique for evaluating view queries is known as

## View Expansion

```
SELECT COUNT(*)  
FROM top_employee  
WHERE department = 'HR'
```

References to a  
view are replaced  
by its definition



```
SELECT COUNT(*)  
FROM (SELECT name  
      FROM employee  
      WHERE salary > 10000)  
WHERE name = 'Joaquim'
```

# Views, data independence and security

Views can *map* data from tables to a new *logical model*.

- **Views** support logical independence from the *physical model*
- **Views** are useful for security context:
  - The DBA can views and grant them access to a group of users.

# Updatable Views

SQL: 1999 distinguishes between two types of views:

- **Updatable Views:** Whose rows can be modified
  - A column of a view can be updated if it is obtained from exactly a **base table** and the **primary key of the base table** is included in the columns of the view
- **Insertable Views:** And views where new lines can be inserted
  - There must be a **one-to-one relationship** between the rows of the view and those of the respective base tables.

# Updatable view: Example

This View is updatable

```
CREATE VIEW senior_employees(eid, emp_name, birthdate)
AS
    SELECT e.eid, e.name as emp_name, e.birthdate
    FROM employee e
    WHERE birthdate < '01-01-1980';
```

```
INSERT INTO senior_employees VALUES (7, 'Grace', '01-12-1979');
SELECT * FROM senior_employees;
```

```
UPDATE senior_employees
SET birthdate = '01-12-1989'
WHERE eid = 7;
```

```
SELECT * FROM senior_employees;
```

# Non-updatable view: Example

This view is **not** updatable!

```
CREATE VIEW senior_managers(emp_name, dep_name, birthdate) AS
  SELECT e.name AS emp_name,
         d.name AS dep_name,
         e.birthdate
  FROM employee e
       JOIN department d ON e.eid = d.mid
 WHERE birthdate < '01-01-1980';
```

```
SELECT * FROM senior_managers;
```

emp_name	dep_name	birthdate
Caroline	Sales	1971-04-07

This statement fails

```
INSERT INTO senior_managers VALUES ('Grace', 'Finance',
'08-03-1979');
```

# Materialised Views

Views sometimes can be too costly to compute and need to be optimised

- **Memory cache** when the query underlying the view is complex and the results fit in memory
  - **Disk Cache** when the query underlying the view is complex and the results do **not** fit in memory
- *How are updates to the underlying table propagated?*



# Materialized Views

Large complex views can be materialised to optimise query performance (over those views)

```
CREATE MATERIALIZED VIEW account_stats AS SELECT ...
```

- *Materialized query tables* and materialised *summary tables* on IBM DB2

```
CREATE TABLE account_stats AS SELECT ...
```

# Refreshing Materialized Views

Also known as *updating views* or *propagating updates*

In POSTGRES, views can be refreshed

- *ON-DEMAND*

```
REFRESH MATERIALIZED VIEW account_stats;
```

- *AUTOMATICALLY*

```
CREATE UNIQUE INDEX idx_account_stats  
ON account_stats(name);
```

```
REFRESH MATERIALIZED VIEW CONCURRENTLY  
account_stats;
```

# Resolving Queries over Materialised Views

Queries over materialised views can be resolved using the *view expansion* strategy or using *query rewriting over views*

```
SELECT name, COUNT(*)  
FROM depositor d  
GROUP BY name  
HAVING count(*) > 2
```



```
SELECT name, num_accts  
FROM account_stats  
WHERE num_accts > 2
```

Detects that we are doing an aggregation query and rewrites the query at the expense of the view (or *existing* materialized views)